




自然語言處理與應用

Natural Language Processing and Applications

Decoding Strategies

Instructor: 林英嘉 (Ying-Jia Lin)
2025/03/31



[Course GitHub](#)



[Slido # NLP_0331](#)

Outline

- Recap: Language Generation
- Decoding Strategies
 - Greedy Decoding
 - Beam Search
 - Top-k / Top-p Sampling
- Evaluations

Natural Language Generation (NLG)

- Natural language generation (NLG) is a **process** that that **outputs** text.
- NLG includes a wide variety of NLP tasks.

Machine
Translation

Abstractive
Summarization

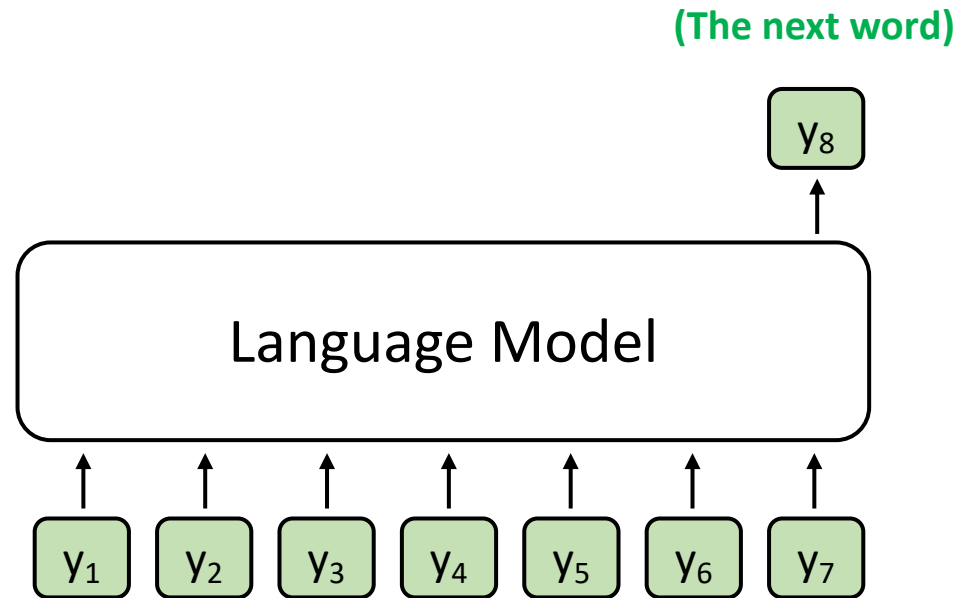
Dialogue
Generation
(e.g., ChatGPT)

Story
Generation

Image
Captioning

...

Recap: Language Model

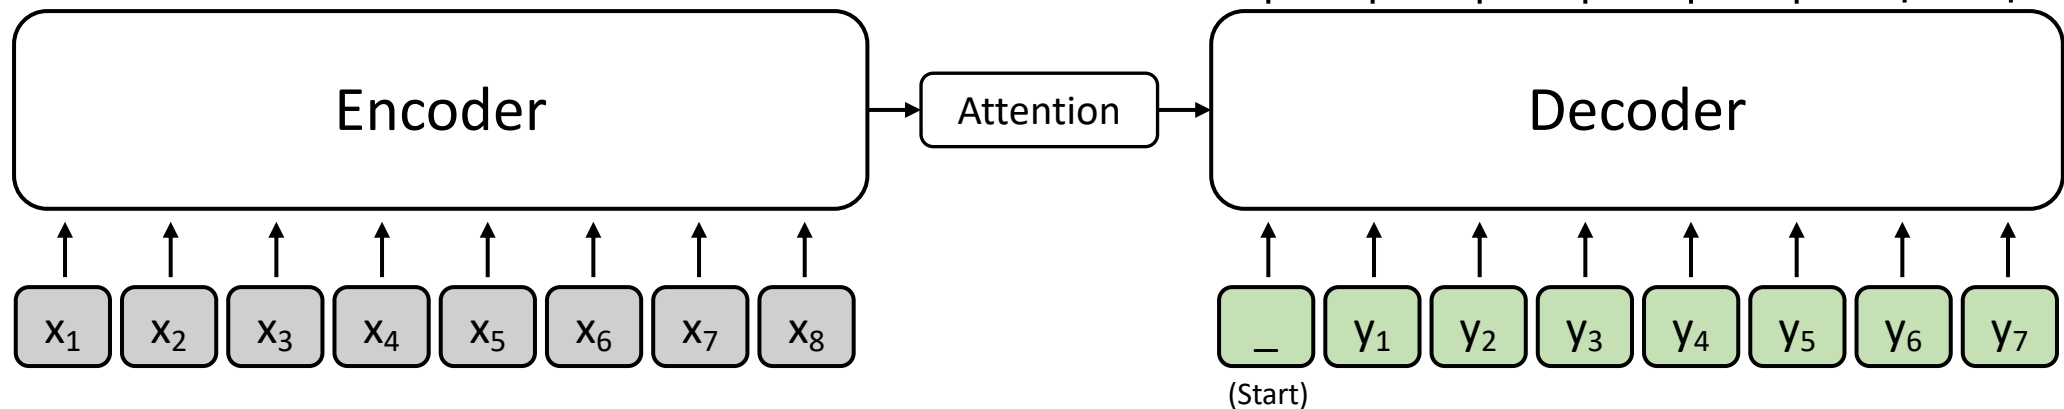


$$P(y_t | y_1, y_2, \dots, y_{t-1})$$

- A model that assigns probabilities to upcoming words is called a **language model**.
- The task involving predictions of upcoming words is **language modeling**.

Recap: Conditional Language Model

- In addition to previous words, a conditional language model is provided with source text x .
- Also referred to sequence-to-sequence models.



Tasks of Conditional Language Model

- In addition to previous words (target), a conditional language model is provided with source text x .

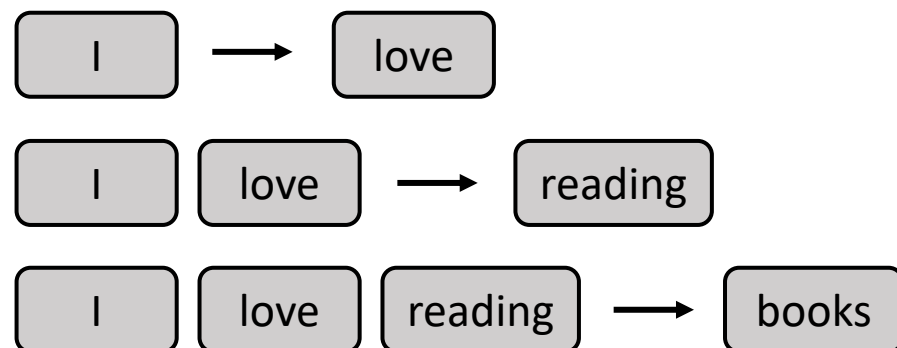
	Source	Target
Machine Translation	Language A	Language B
Summarization	Long Text	Concise Text
Dialogue Generation	User Input	Desired User Input
...		

How to train a (Conditional) Language Model?


- First, you need a training corpus.

Example: I love reading books.

Language modeling (**Unsupervised**)

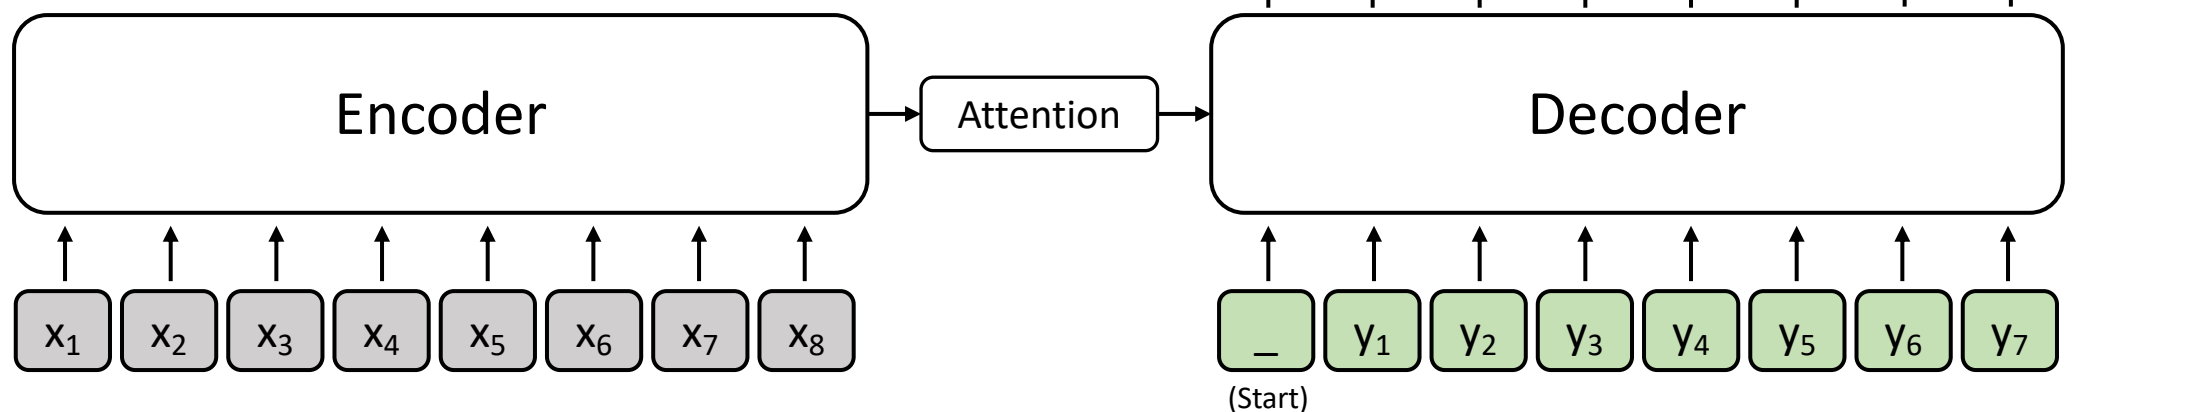


Machine Translation (**Supervised**)



How to train a (Conditional) Language Model?

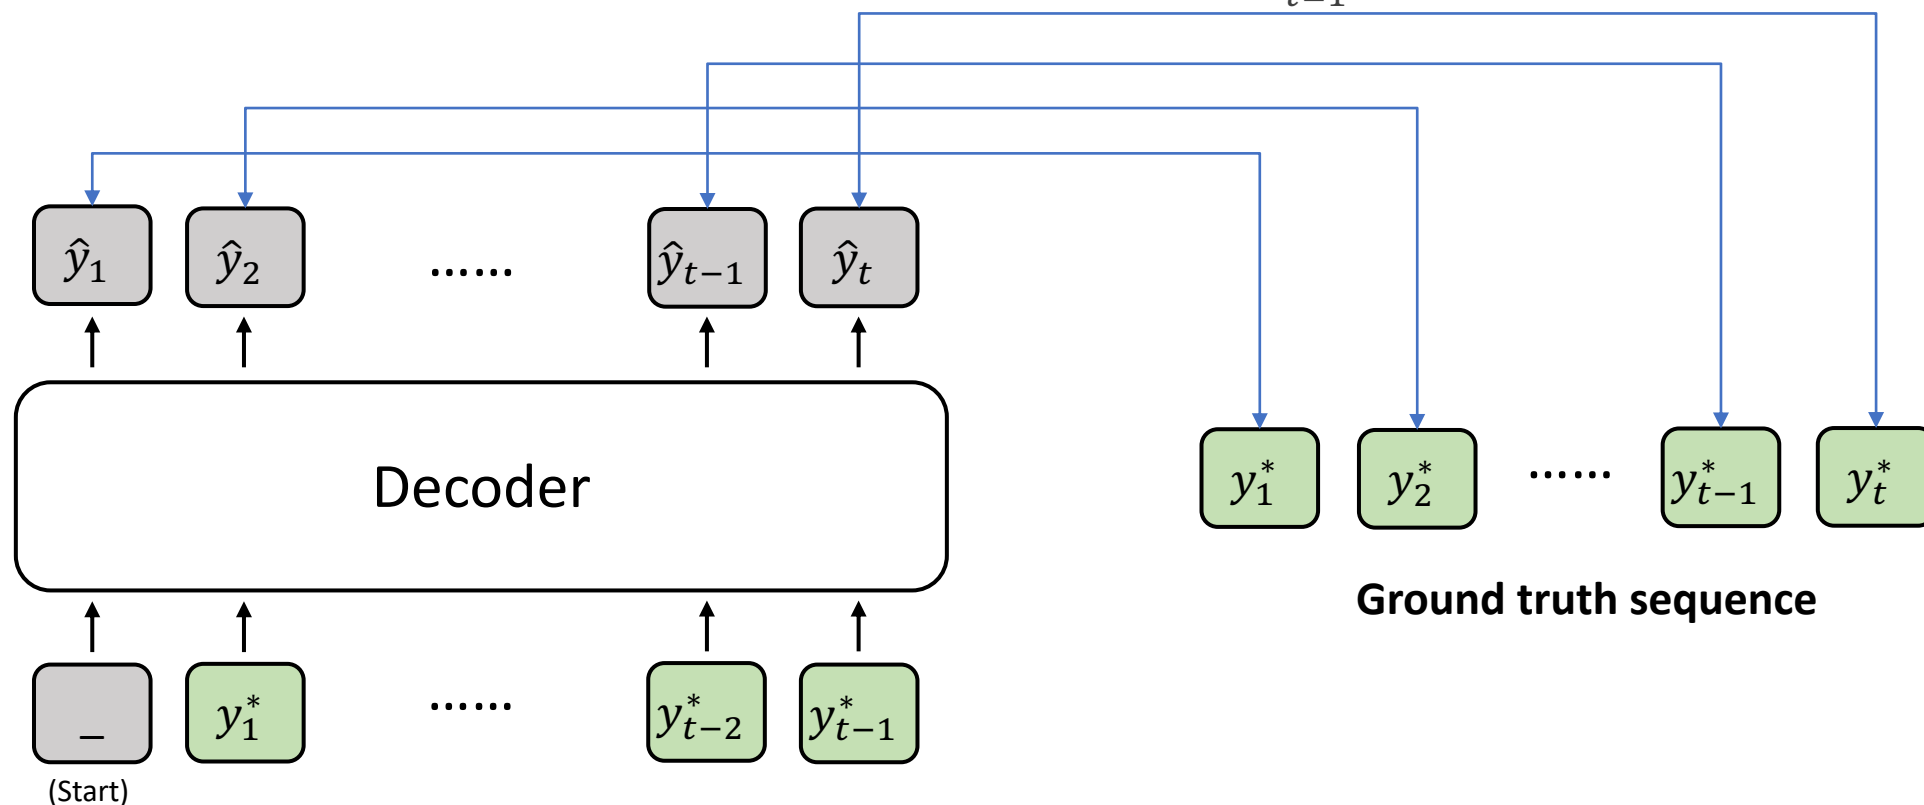
- Use the Teacher Forcing technique during training.
- Total loss for a sequence: $\sum_1^T l_t$
 - T : Sequence length



Teacher Forcing – Training Time


During training:

$$L_{ml} = - \sum_{t=1}^{n'} \log p(y_t^* | y_1^*, \dots, y_{t-1}^*, x)$$



Teacher Forcing – Testing Time

During testing:



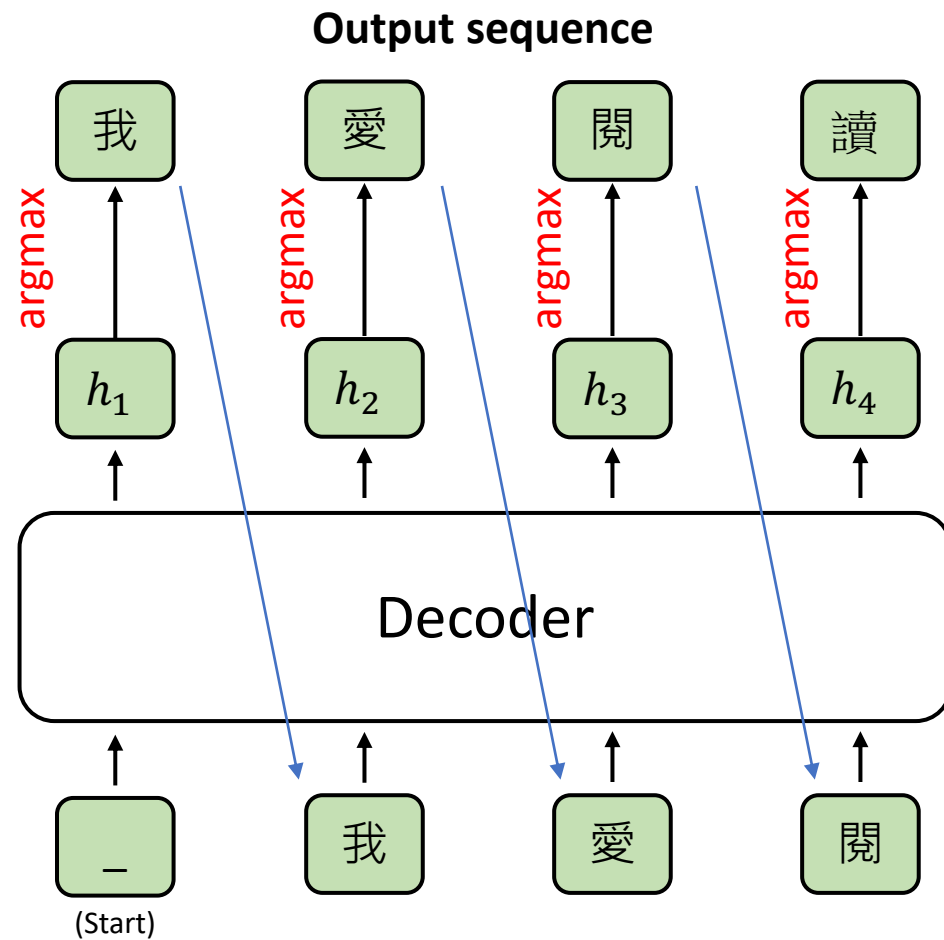
- Advantage: stabilize training and increase performance
- Question: **How does the next word be determined?**

Decoding Strategies

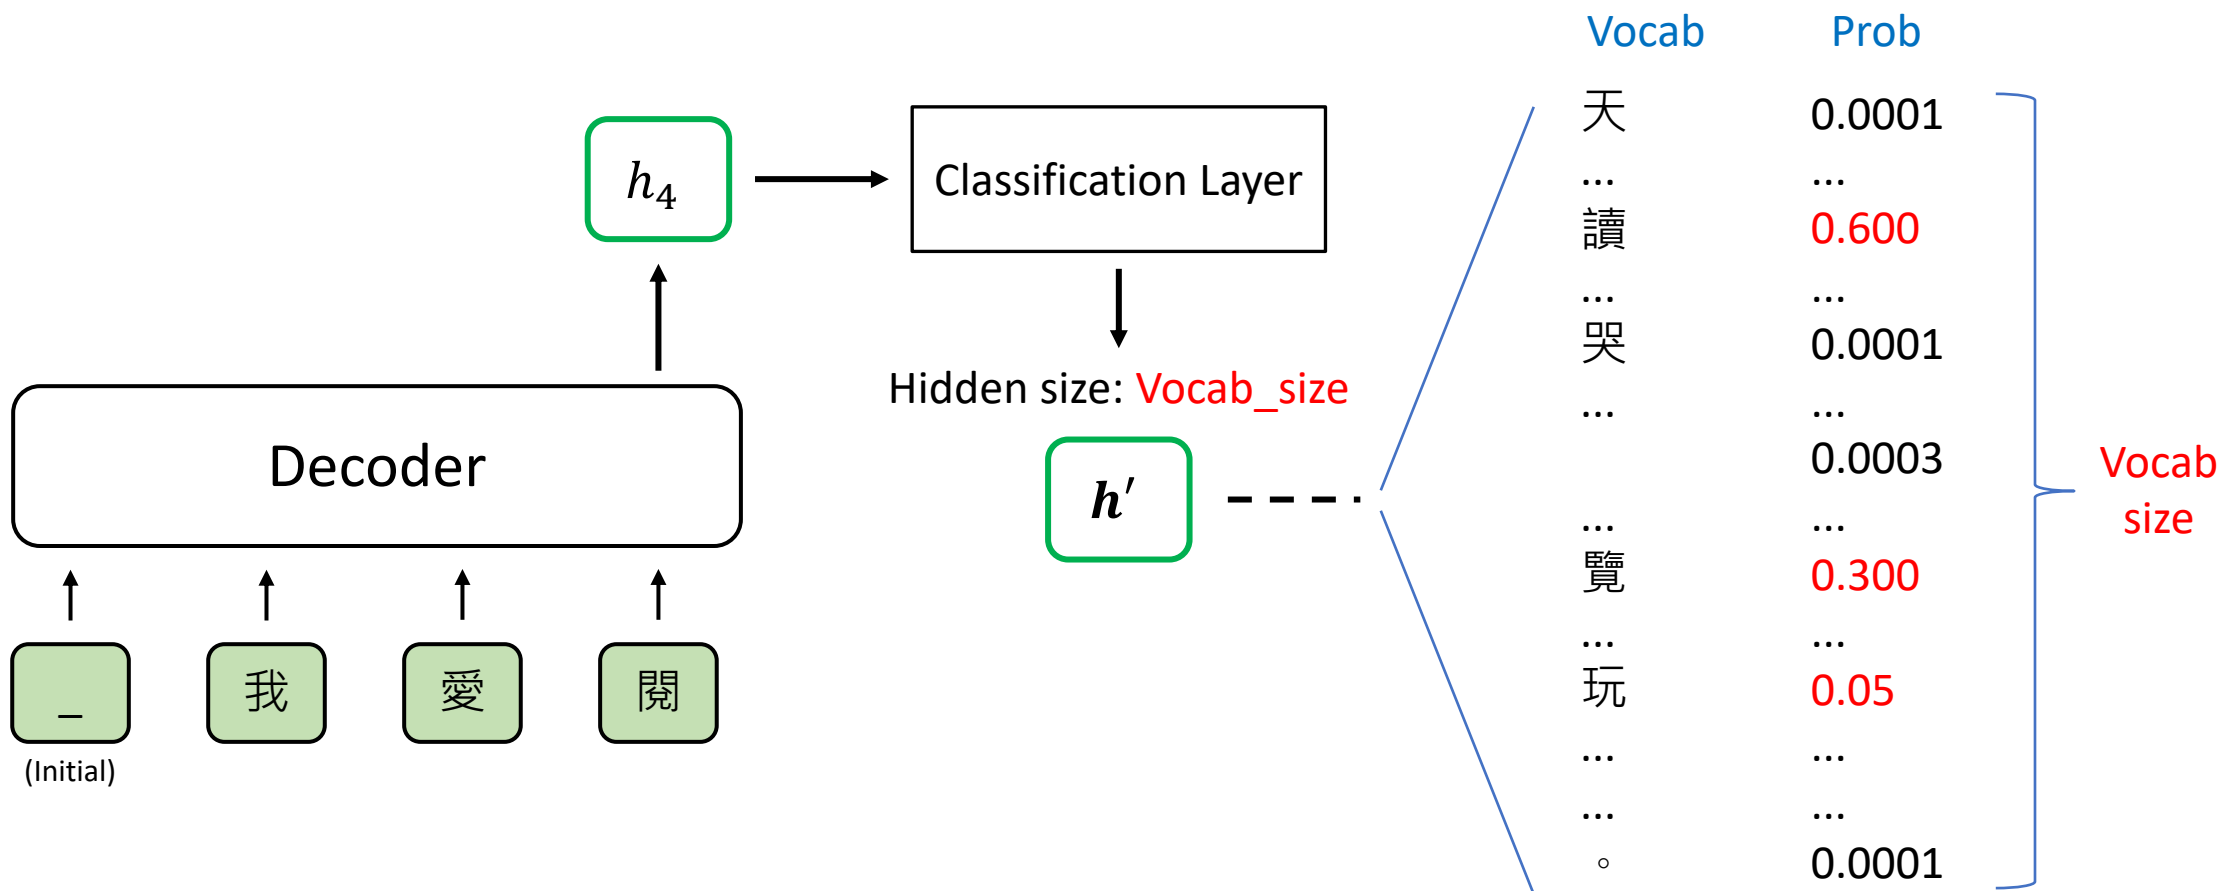
- Greedy Decoding
- Beam Search
- Top-k Sampling
- Top-p Sampling

Greedy Decoding

Example: I love reading books.

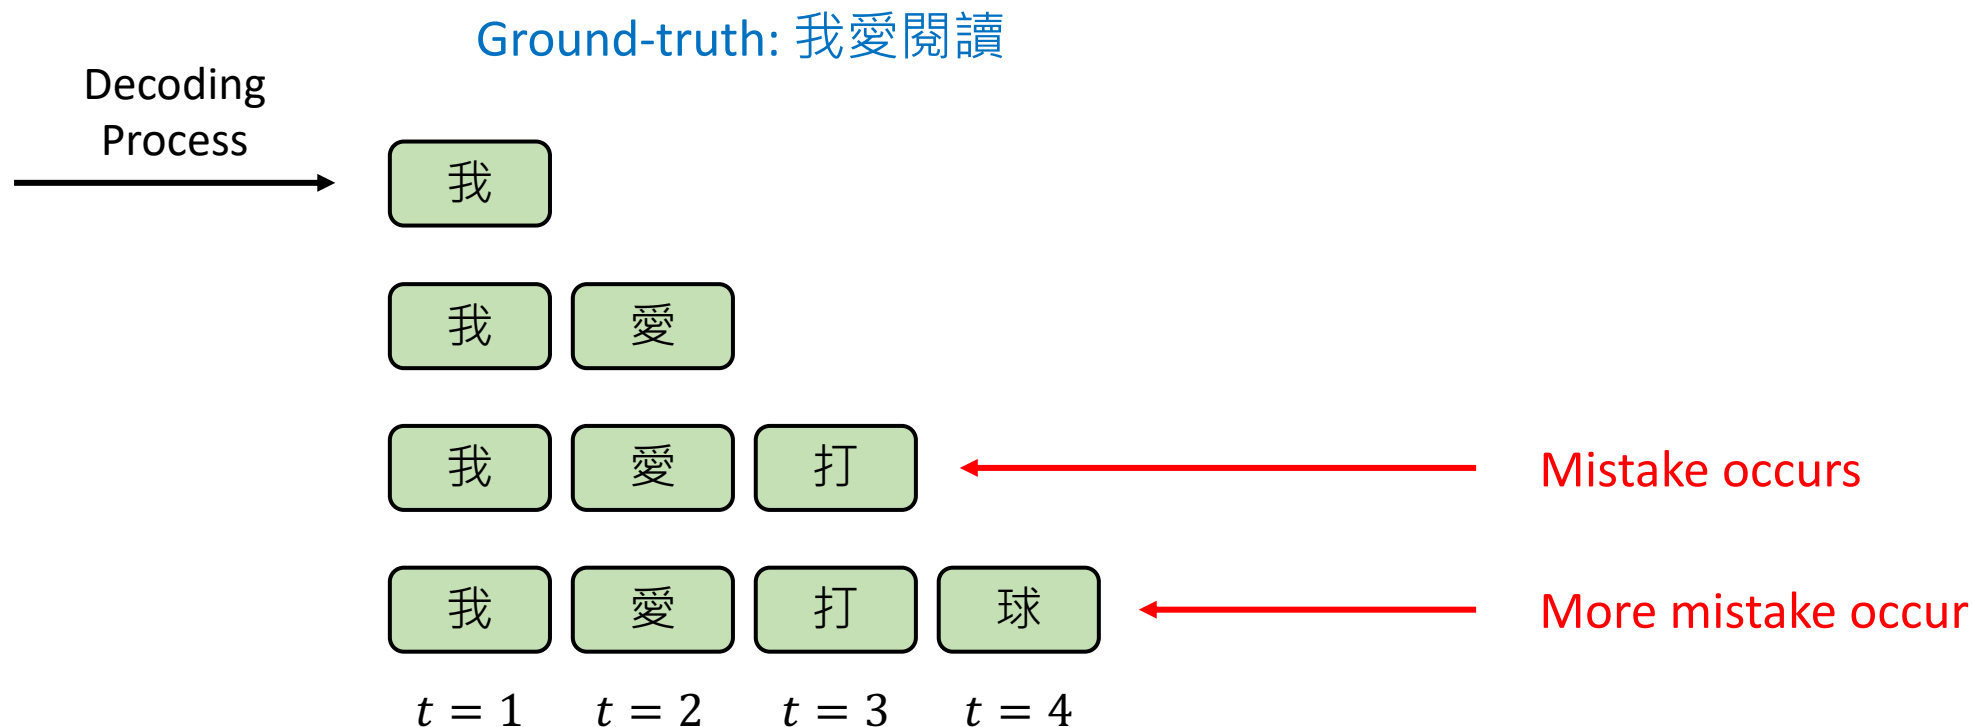


Greedy Decoding – Best Selection Process



Problem of Greedy Decoding

- Greedy decoding cannot undo!



Re-thinking Greedy Decoding

- Greedy decoding cannot undo!
- Greedy decoding only provides one best choice at each time step.
- How about providing **more than one choices** at each time step?



Beam Search


Beam Search

- Set the `Beam size` (or `Beam width`) = 2
 - This means that the number of candidates will be preserved at each decoding time.
 - Beam size is a hyperparameter for beam search decoding.
- At each decoding time step, a score is calculated via the following equation:

$$L_{ml} = \sum_{t=1}^{n'} \log p(y_t^* | y_1^*, \dots, y_{t-1}^*, x)$$

Beam Search ($t = 1$)


`Beam size` = 2



- At this decoding step, two choices are preserved.

Beam Search ($t = 1$)

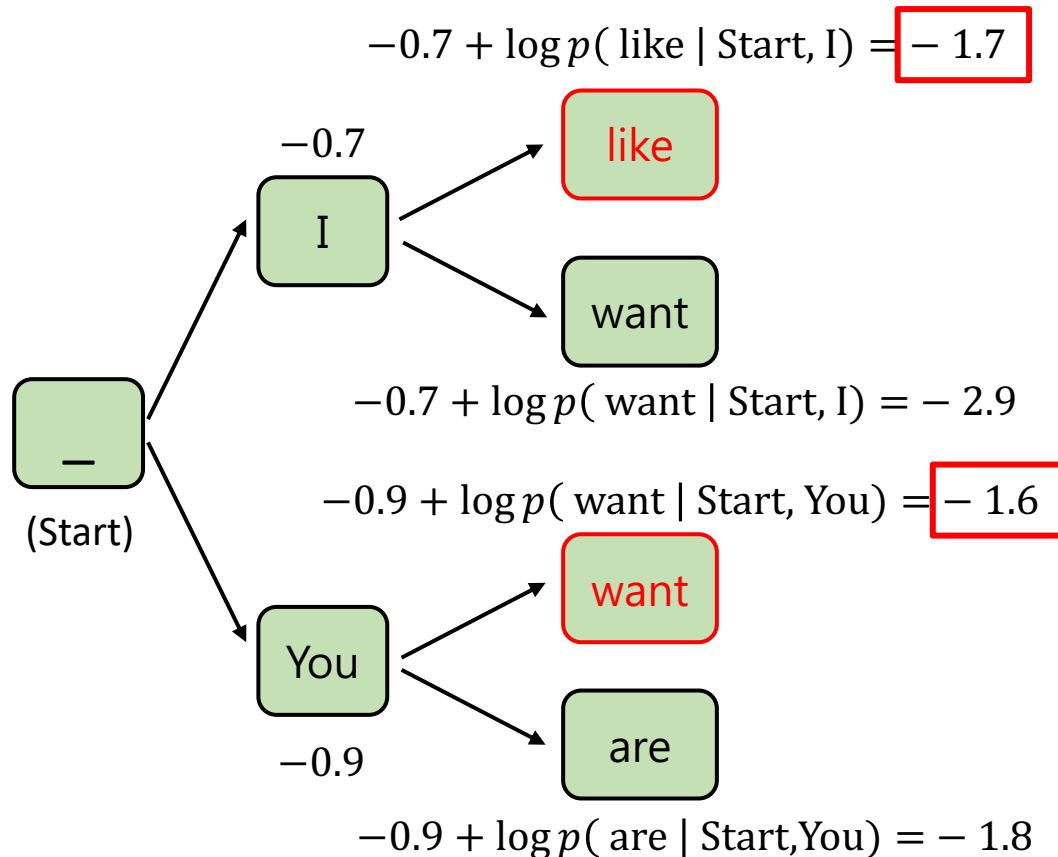
`Beam size` = 2



- At this decoding step, two choices are preserved.

Beam Search ($t = 2$)

Beam size = 2




$$L_{ml} = - \sum_{t=1}^T \log p(y_t^* | y_1^*, \dots, y_{t-1}^*, x)$$

Note the loglikelihood! Being close to zero is better!

- At this decoding step, two choices are preserved, and the other two are discarded.

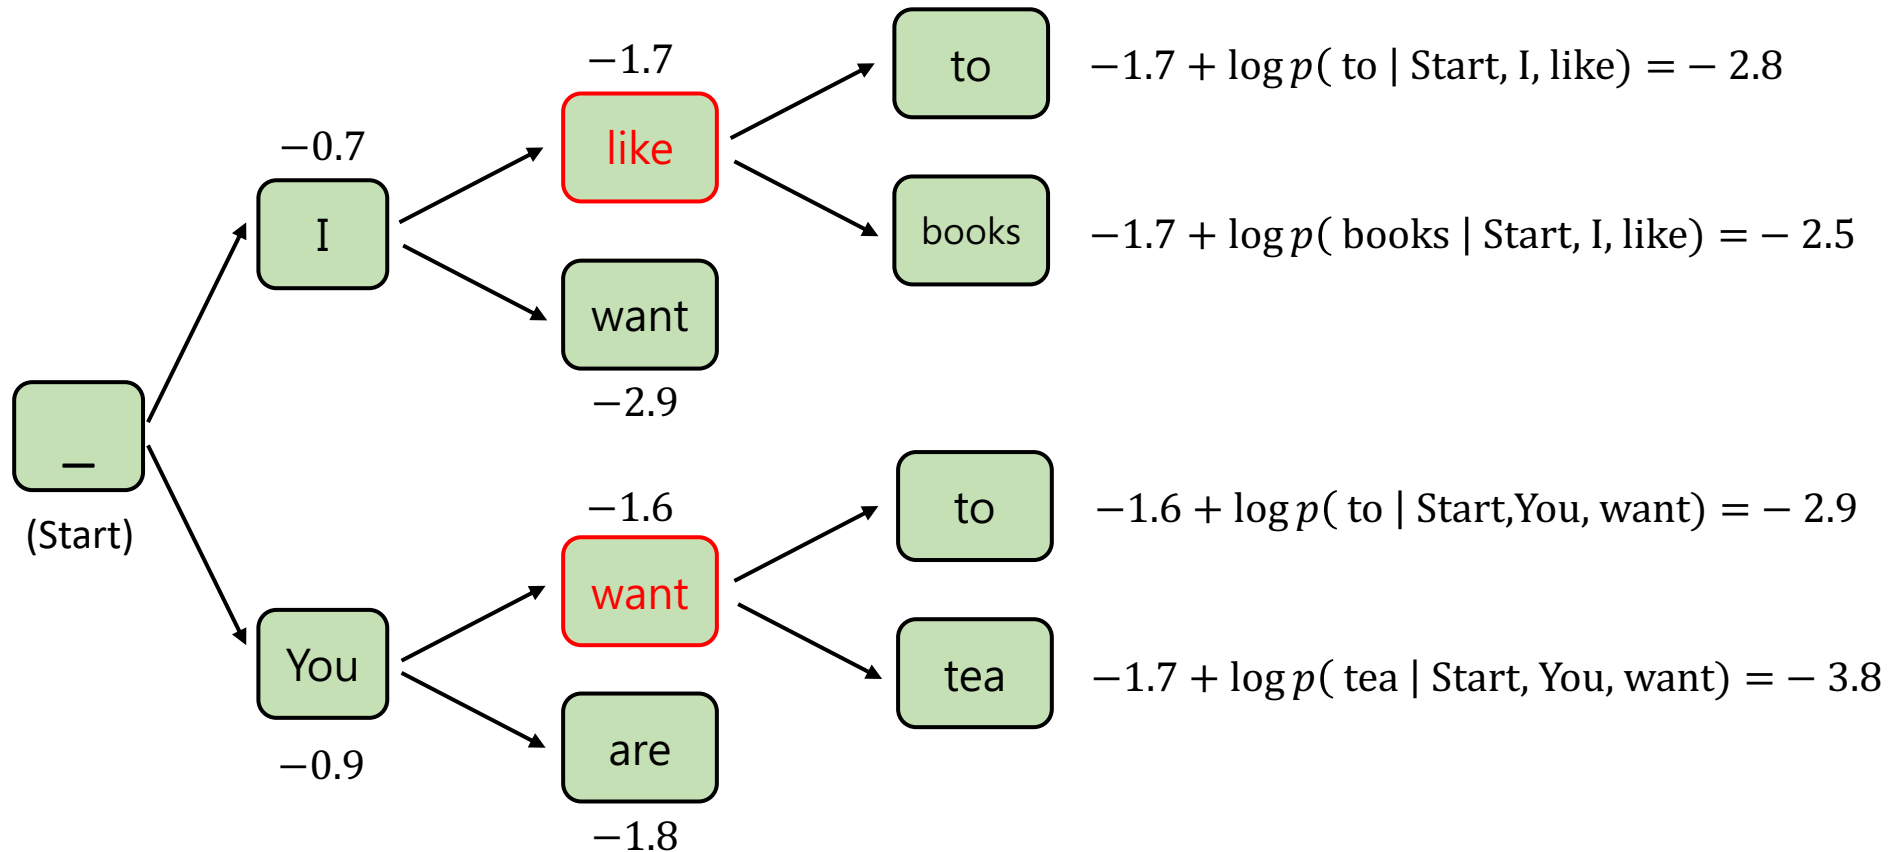
Beam Search ($t = 2$)

`Beam size` = 2



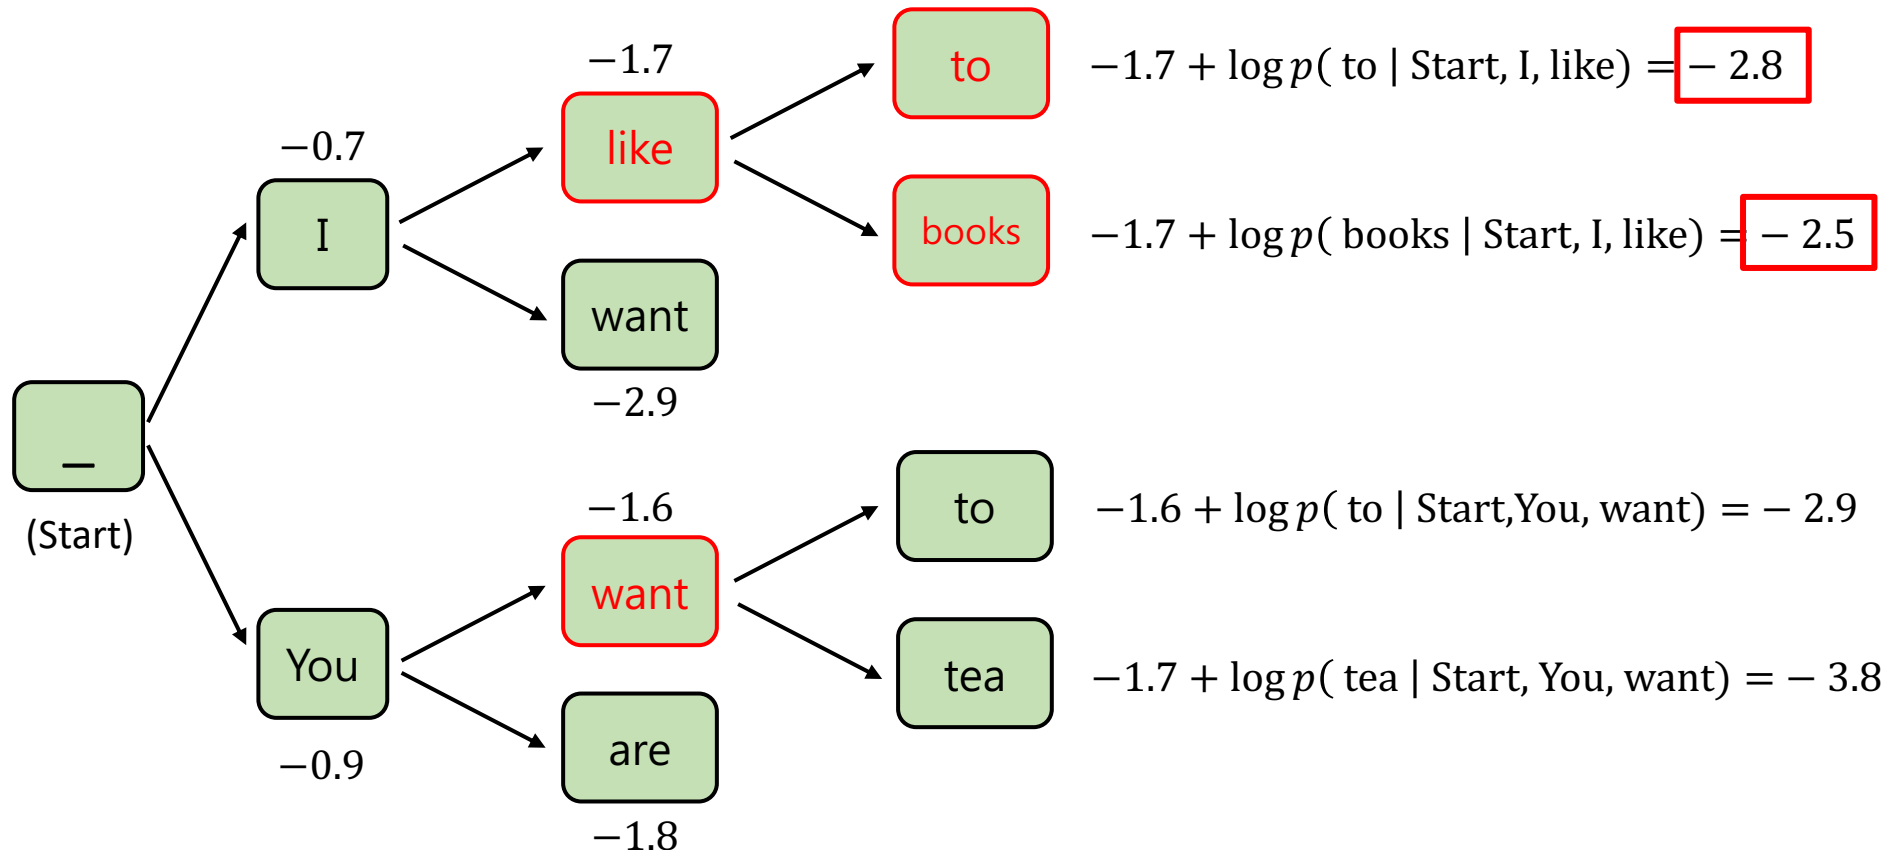
Beam Search ($t = 3$)

Beam size = 2




Beam Search ($t = 3$)

'Beam size' = 2



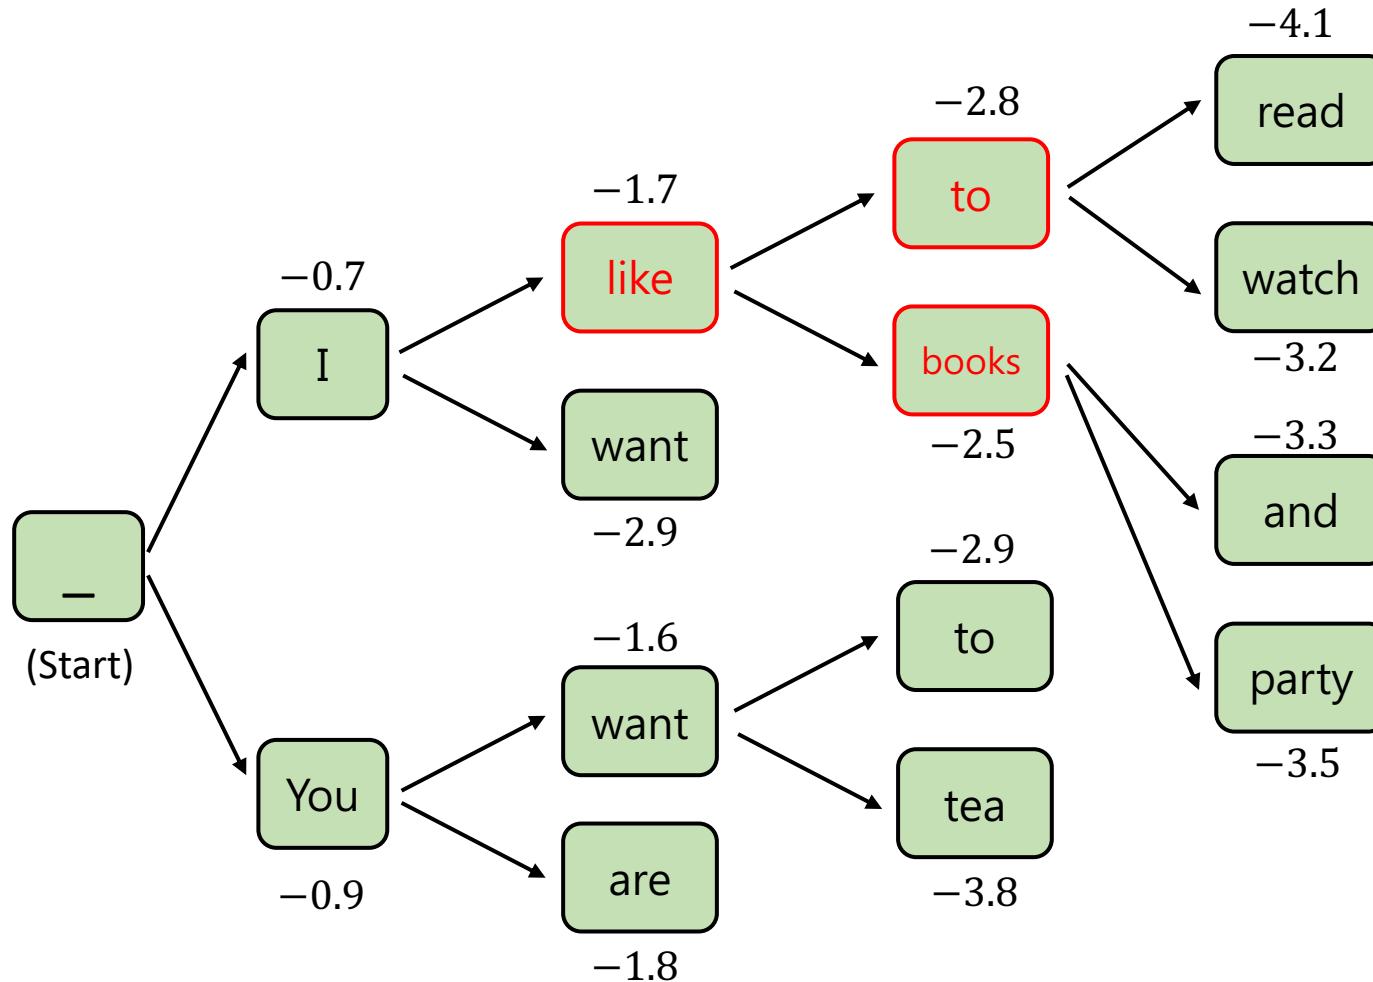
Beam Search ($t = 3$)

`Beam size` = 2



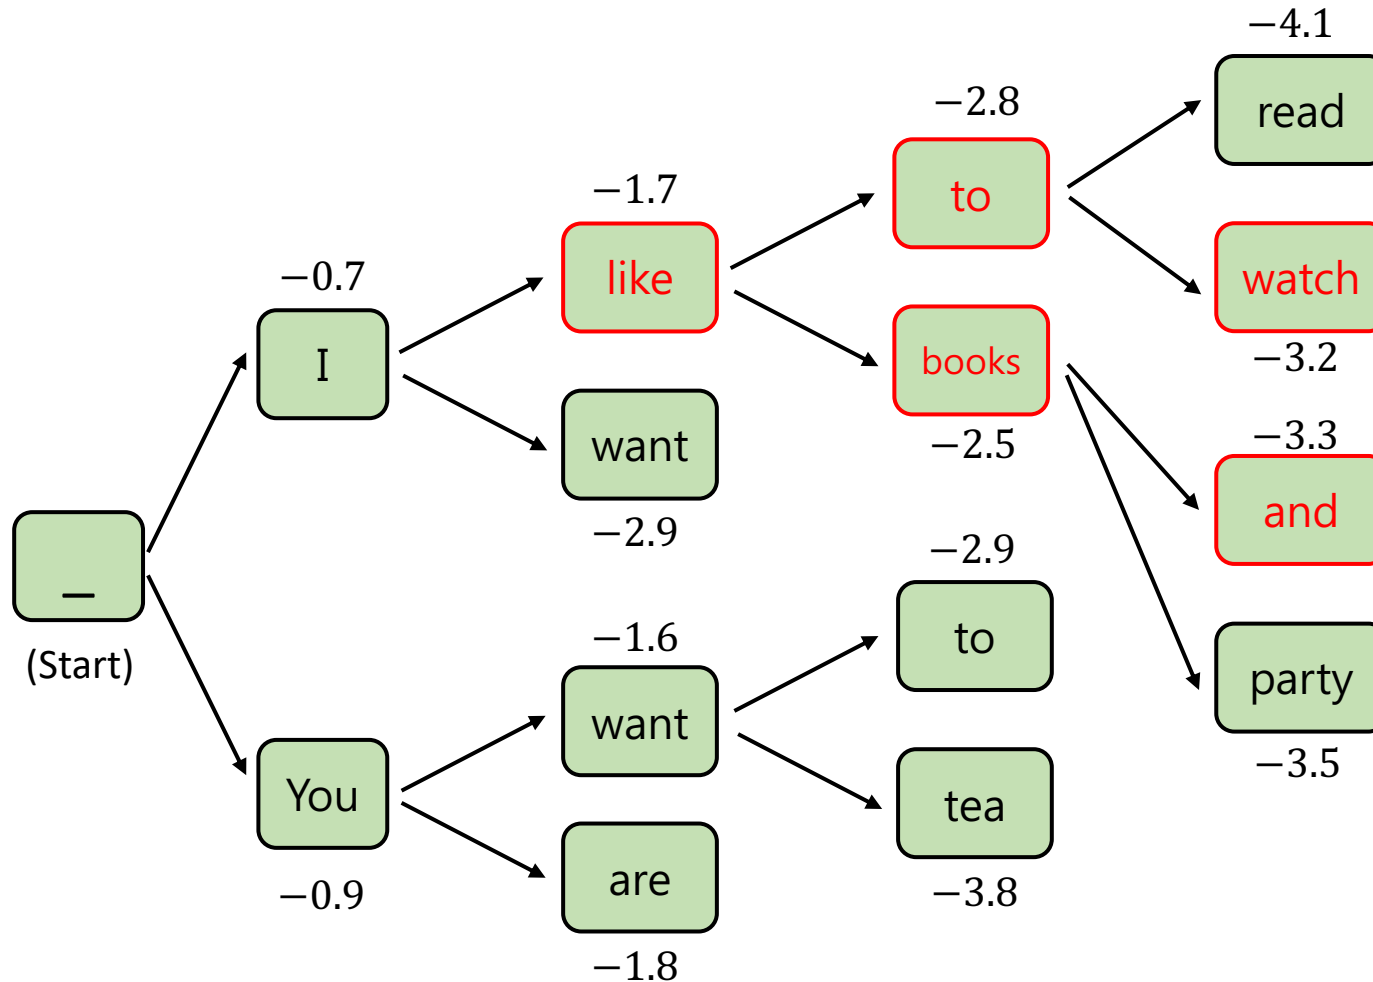
Beam Search ($t = 4$)

`Beam size` = 2




Beam Search ($t = 4$)

`Beam size` = 2




Beam Search ($t = 5$)

'Beam size' = 2




Beam Search ($t = 5$)

'Beam size' = 2



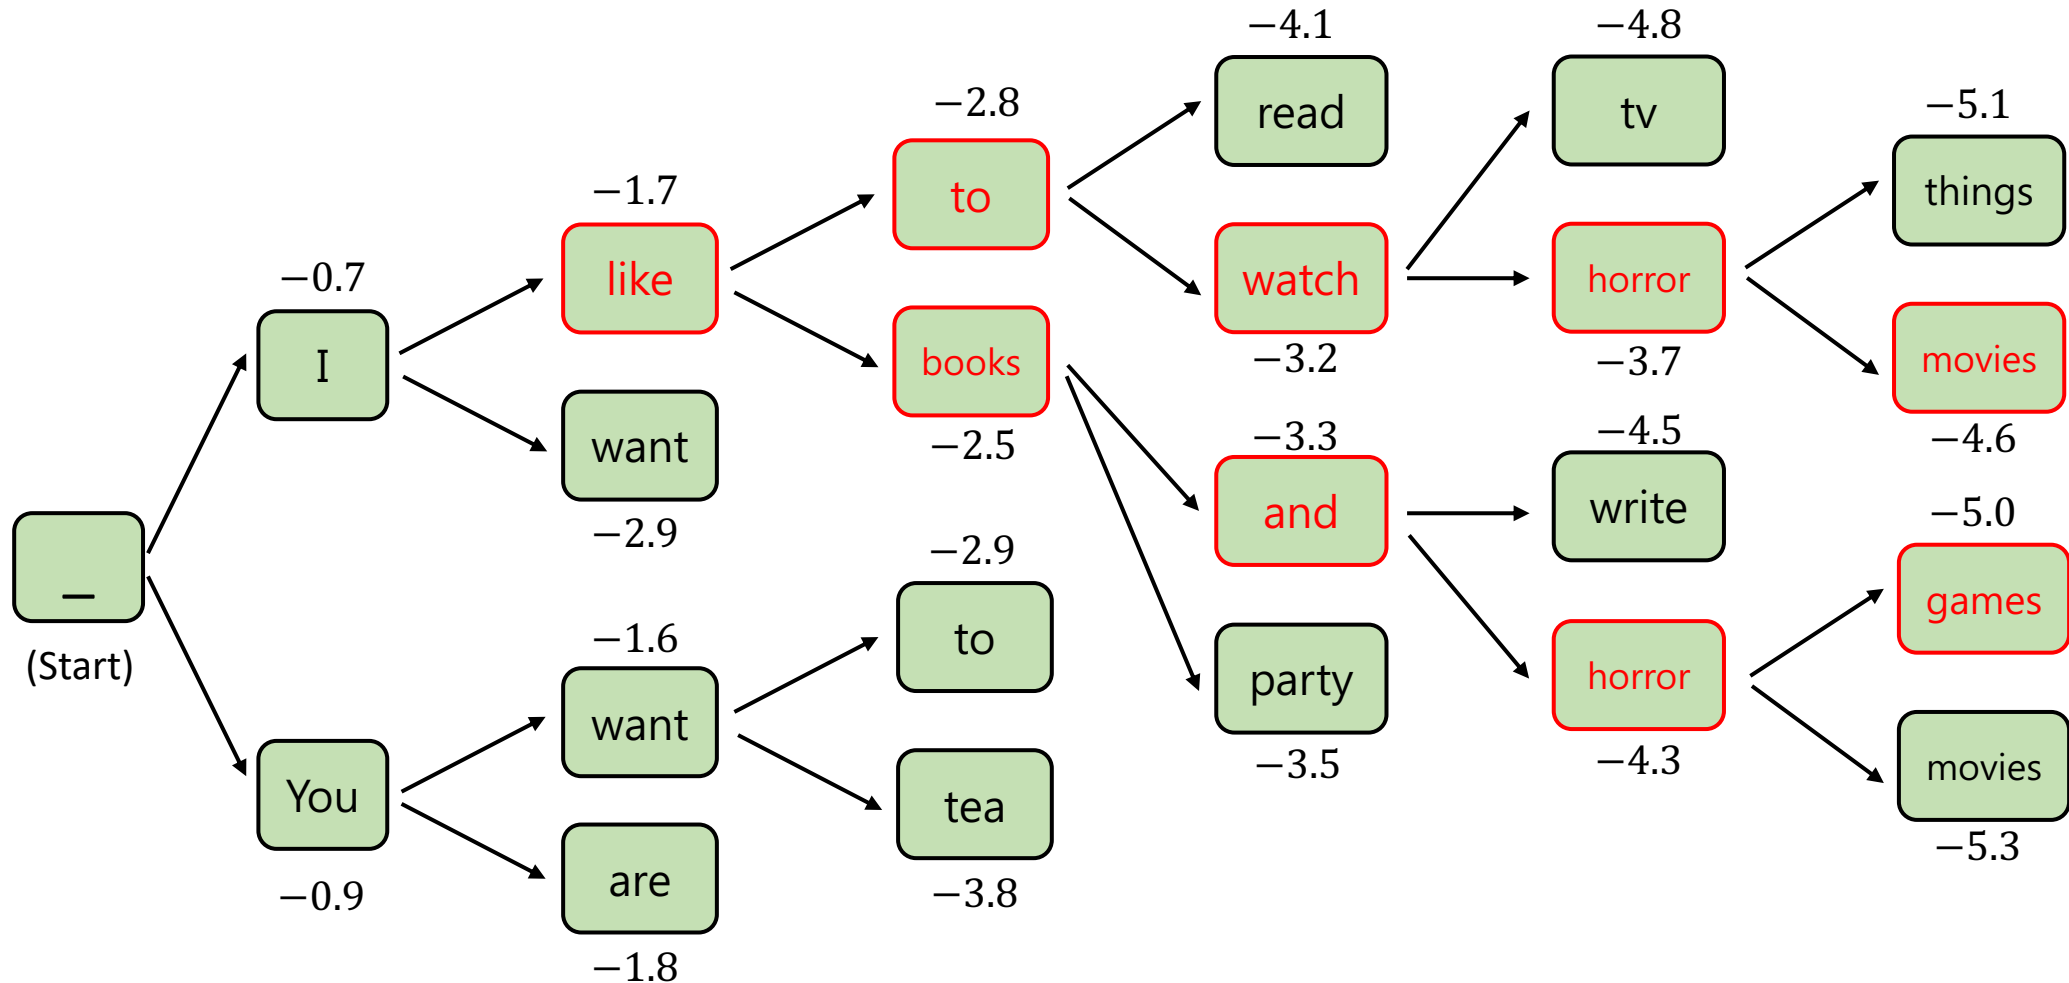
Beam Search ($t = 6$)

'Beam size' = 2



Beam Search ($t = 6$)

'Beam size' = 2



Stop Criterion


- There are two common stop criterion, either for greedy decoding or beam search decoding:
 - We consider a sequence of generation complete when the <EOS> token is produced by a model. *<EOS>: End of sequence
 - E.g., <Start> I like to watch horror movies <EOS>
- A generated sequence reaches a pre-defined **maximal length**.

Problem of Beam Search

- Longer candidates will have lower scores.
- (Let's see again the 6th time step)

Beam Search ($t = 6$)

'Beam size' = 2



Problem of Beam Search (1)

- Longer candidates will have lower scores.
- Solution: Perform normalization to penalize on length

$$L_{ml} = -\frac{1}{T} \sum_{t=1}^T \log p(y_t^* | y_1^*, \dots, y_{t-1}^*, x)$$

Problem of Beam Search (2)

- Decoding strategies that optimize for output with high probability, such as beam search, lead to text that is **incredibly degenerate (e.g., repetitive words)**, even when using state-of-the-art models such as GPT-2 Large (in 2020).

Holtzman, Ari, et al. "The Curious Case of Neural Text Degeneration." International Conference on Learning Representations (ICLR). 2020.

Problem of Beam Search (2) – Continued.

Context: In a shocking finding, scientist discovered a herd of unicorns living in a remote, previously unexplored valley, in the Andes Mountains. Even more surprising to the researchers was the fact that the unicorns spoke perfect English.

Beam Search, $b=32$:

"The study, published in the Proceedings of the National Academy of Sciences of the United States of America (PNAS), was conducted by researchers from the Universidad Nacional Autónoma de México (UNAM) and the Universidad Nacional Autónoma de México (UNAM/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de ..."

Holtzman, Ari, et al. "The Curious Case of Neural Text Degeneration." International Conference on Learning Representations. 2019.

Problem of Beam Search (2) – Continued.




Figure 4: The probability of a repeated phrase increases with each repetition, creating a positive feedback loop. We found this effect to hold for the vast majority of phrases we tested, regardless of phrase length or if the phrases were sampled randomly rather than taken from human text.

Holtzman, Ari, et al. "The Curious Case of Neural Text Degeneration." International Conference on Learning Representations. 2019.

Problem of Beam Search (2) – Continued.

Context: In a shocking finding, scientist discovered a herd of unicorns living in a remote, previously unexplored valley, in the Andes Mountains. Even more surprising to the researchers was the fact that the unicorns spoke perfect English.

Beam Search, $b=32$:

"The study, published in the Proceedings of the National Academy of Sciences of the United States of America (PNAS), was conducted by researchers from the Universidad Nacional Autónoma de México (UNAM) and the Universidad Nacional Autónoma de México (UNAM/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de ..."

How about random sampling
according to the probabilities?
(Pure Sampling)

Holtzman, Ari, et al. "The Curious Case of Neural Text Degeneration." International Conference on Learning Representations. 2019.

Problem of Beam Search (2) – Continued.

Context: In a shocking finding, scientist discovered a herd of unicorns living in a remote, previously unexplored valley, in the Andes Mountains. Even more surprising to the researchers was the fact that the unicorns spoke perfect English.

Beam Search, $b=32$:

"The study, published in the Proceedings of the National Academy of Sciences of the United States of America (PNAS), was conducted by researchers from the Universidad Nacional Autónoma de México (UNAM) and the Universidad Nacional Autónoma de México (UNAM/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de ..."

Pure Sampling:

They were cattle called Bolivian Cavalleros; they live in a remote desert uninterrupted by town, and they speak huge, beautiful, paradisiacal Bolivian linguistic thing. They say, 'Lunch, marge.' They don't tell what the lunch is," director Professor Chuperas Omwell told Sky News. "They've only been talking to scientists, like we're being interviewed by TV reporters. We don't even stick around to be interviewed by TV reporters. Maybe that's how they figured out that they're cosplaying as the Bolivian Cavalleros."

Holtzman, Ari, et al. "The Curious Case of Neural Text Degeneration." International Conference on Learning Representations. 2019.

Observation between BS and Pure Sampling

- Pure Sampling does not show repetitive loop, but the result becomes incoherent and almost unrelated to the context
- Why? -> **Unreliable tail**

Words that have low probabilities

Vocab	Prob	Example Vocab
the	0.0011	
am	0.0012	
no	0.0013	
a	0.0014	
/	0.0015	
...	...	
...	...	
...	...	
...	...	
...	...	
...	...	
...	...	
...	...	

Observation between BS and Pure Sampling

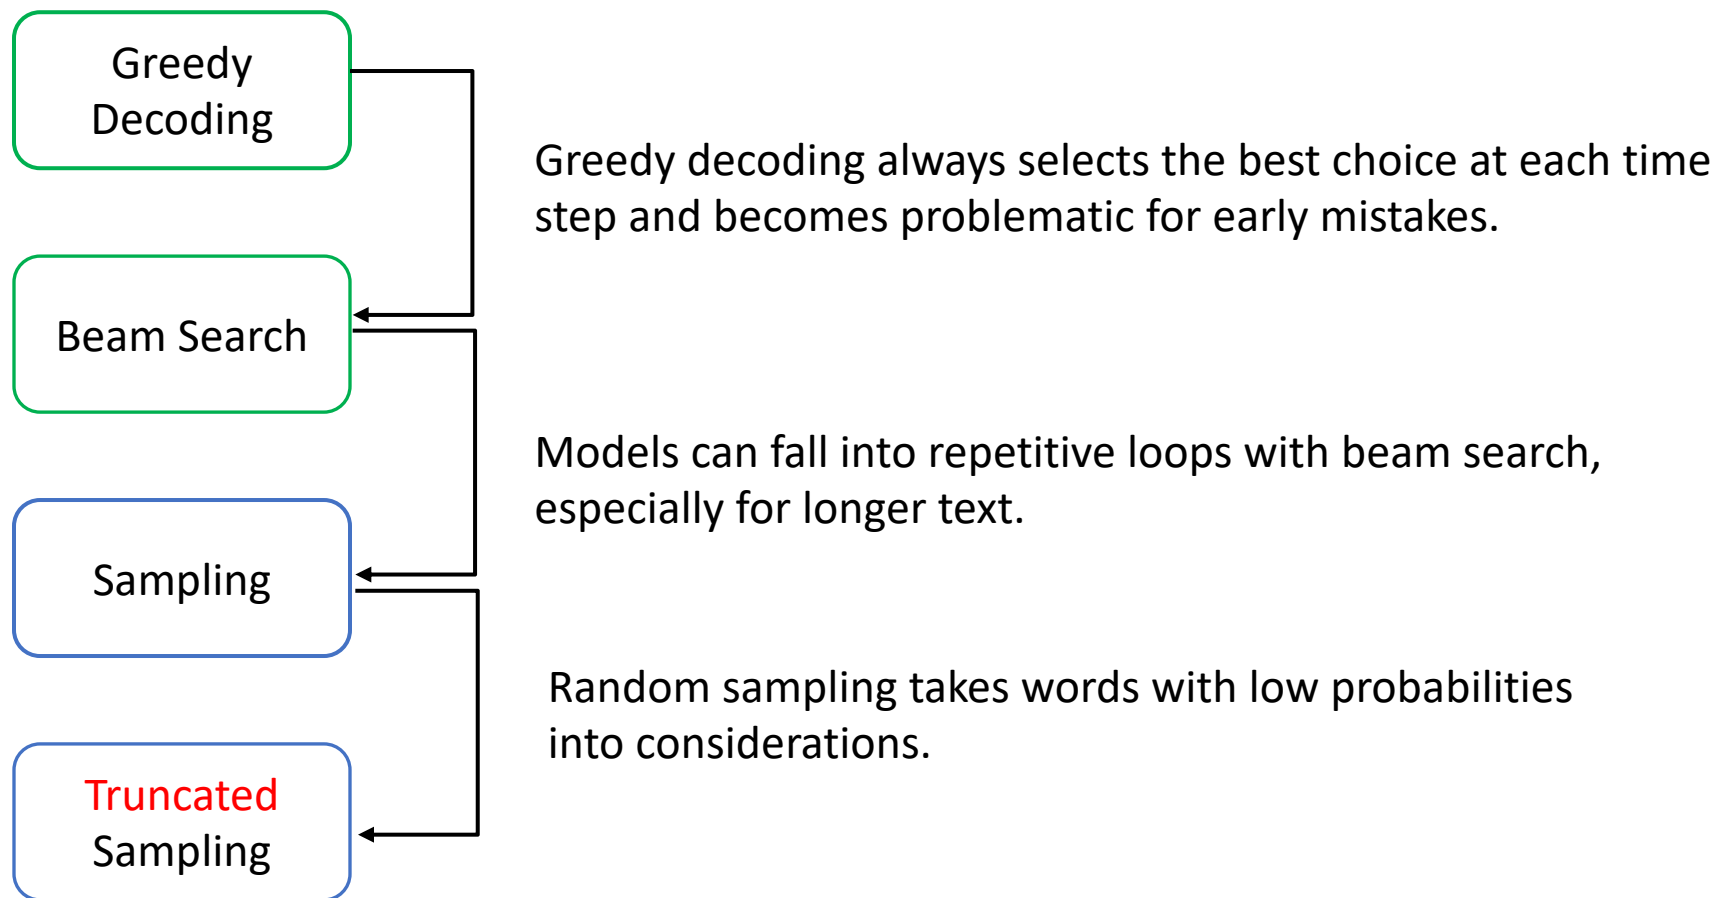
- Pure Sampling does not show repetitive loop, but the result becomes incoherent and almost unrelated to the context
- Why? -> Unreliable tail

Words that have low probabilities

Vocab	Prob	Example Vocab
the	0.0011	
am	0.0012	
no	0.0013	
a	0.0014	
/	0.0015	
...	...	
...	...	
...	...	
...	...	
...	...	

Let's truncate the vocabulary!

Summary and the Thinking Route



Top-k Sampling

*To prevent confusion, we call each unit in a vocabulary as “word” instead of “token”, even though it is a sub-word unit.

- Core idea: truncate the vocabulary with **the most probable words**
- Steps:
 1. Define a value k as the size of truncated vocabulary.
 2. Leave the k words with the highest probabilities. Now you get a new truncated vocabulary $V^{(k)}$.
 3. Re-build the probability distribution based on the following equations:

$$p' = \sum_{x \in V^{(k)}} P(x|x_{1:x-1})$$

$$P'(x|x_{1:x-1}) = \begin{cases} P(x|x_{1:x-1}/p') & \text{if } x \in V^{(k)} \\ 0 & \text{otherwise} \end{cases}$$

Top-p Sampling

- Core idea: truncate the vocabulary based on **probability mass**
- Steps:
 1. Define a value p as the probability threshold.
 2. Truncate the vocabulary according to the following equation:

$$\sum_{x \in V^{(p)}} P(x|x_1:x_{-1}) \geq p$$

where $V^{(p)} \subset V$ is the smallest set that fulfills the equation.
Now you get a new truncated vocabulary $V^{(p)}$.

Top-p Sampling

- Core idea: truncate the vocabulary based on **probability mass**
- Steps:

3. Re-build the probability distribution based on the following equations:

$$p' = \sum_{x \in V(p)} P(x|x_{1:x-1})$$

$$P'(x|x_{1:x-1}) = \begin{cases} P(x|x_{1:x-1}/p') & \text{if } x \in V(p) \\ 0 & \text{otherwise} \end{cases}$$

Softmax

- When generating the next word, `softmax` is performed to get the probabilities among the words in the vocabulary.
- Softmax formula:

$$\frac{\exp(u_l/t)}{\sum_{l'}^{|V|} \exp(u_{l'}/t)}$$

u_l : logits (model outputs before softmax)

$|V|$: size of the vocabulary

t : softmax temperature

Softmax

	Word	Logits		Probability
Example Vocab	the	0.0011	————→	0.78
	am	0.0012	————→	0.11
	no	0.0013	————→	0.03
	a	0.0014	————→	0.02
	/	0.0015	————→	0.06

- Note that softmax is required for every decoding strategies since we need to find out the next word from a vocabulary.

Properties of Softmax Temperature

- Softmax formula:

$$\frac{\exp(u_l/t)}{\sum_{l'}^{|V|} \exp(u_{l'}/t)}$$

u_l : logits (model outputs before softmax)
 $|V|$: size of the vocabulary
 t : softmax temperature

- Larger t -> Lower probability value -> Smaller ranges of probability distribution -> **More diverse** Outputs
- Smaller t -> Higher probability value -> Greater ranges of probability distribution -> **Less diverse** Outputs

Top-k and Top-p Sampling Improve Repetition

Method	Perplexity	Self-BLEU4	Zipf Coefficient	Repetition %	HUSE
Human	12.38	0.31	0.93	0.28	-
Greedy	1.50	0.50	1.00	73.66	-
Beam, b=16	1.48	0.44	0.94	28.94	-
Stochastic Beam, b=16	19.20	0.28	0.91	0.32	-
Pure Sampling	22.73	0.28	0.93	0.22	0.67
Sampling, $t=0.9$	10.25	0.35	0.96	0.66	0.79
Top- $k=40$	6.88	0.39	0.96	0.78	0.19
Top- $k=640$	13.82	0.32	0.96	0.28	0.94
Top- $k=40$, $t=0.7$	3.48	0.44	1.00	8.86	0.08
Nucleus $p=0.95$	13.13	0.32	0.95	0.36	0.97

Table 1: Main results for comparing all decoding methods with selected parameters of each method. The numbers *closest to human scores* are in **bold** except for HUSE (Hashimoto et al., 2019), a combined human and statistical evaluation, where the highest (best) value is **bolded**. For Top- k and Nucleus Sampling, HUSE is computed with interpolation rather than truncation (see §6.1).

Holtzman, Ari, et al. "The Curious Case of Neural Text Degeneration." International Conference on Learning Representations. 2019.

The k value should be carefully picked

Method	Perplexity	Self-BLEU4	Zipf Coefficient	Repetition %	HUSE
Human	12.38	0.31	0.93	0.28	-
Greedy	1.50	0.50	1.00	73.66	-
Beam, b=16	1.48	0.44	0.94	28.94	-
Stochastic Beam, b=16	19.20	0.28	0.91	0.32	-
Pure Sampling	22.73	0.28	0.93	0.22	0.67
Sampling, $t=0.9$	10.25	0.35	0.96	0.66	0.79
Top- $k=40$	6.88	0.39	0.96	0.78	0.19
Top- $k=640$	13.82	0.32	0.96	0.28	0.94
Top- $k=40$, $t=0.7$	3.48	0.44	1.00	8.86	0.08
Nucleus $p=0.95$	13.13	0.32	0.95	0.36	0.97

Table 1: Main results for comparing all decoding methods with selected parameters of each method. The numbers *closest to human scores* are in **bold** except for HUSE (Hashimoto et al., 2019), a combined human and statistical evaluation, where the highest (best) value is **bolded**. For Top- k and Nucleus Sampling, HUSE is computed with interpolation rather than truncation (see §6.1).

Holtzman, Ari, et al. "The Curious Case of Neural Text Degeneration." International Conference on Learning Representations. 2019.

Comparison for Top-k and Top-p Sampling

	Top-k Sampling	Top-p Sampling
Hyperparameter	k: top-k words are preserved	p: sum of the minimum set of words exceeds the value of p
Performance (Who is better?)	By cases (these two are both widely used)	
Hyperparameter Tuning	Harder	Easier
Common Hyperparameter Value	k=40	p=0.95


Evaluations

- Perplexity
- BLEU Score
- ROUGE Score
- BERTScore
- BLEURT

How to evaluate natural language generation?

- Natural language is hard to evaluate due to subjectivity and language diversity.

For example: Machine Translation



- Human evaluations
- Automatic evaluations (We will focus on this topic.)

(Recap) Perplexity

Perplexity (PPL) is a quantitative criterion used to evaluate the capacities of language modeling models. -> **How confident is a model to its output?**

- Given the sequence of words $W = w_1w_2 \dots w_N$ and an N-gram model. The PPL of the model was computed by:

$$\text{Perplexity}(W) = P(w_1w_2 \dots w_N)^{-\frac{1}{N}} = \sqrt[N]{\prod_{k=1}^N \frac{1}{P(w_k|w_{k-N+1:k})}}$$

The lower the value of perplexity, the better the language modeling capability of the model.

BLEU (Bilingual Evaluation Understudy)

- A word-based metric.
 - It is very sensitive to word tokenization
- Core concept: Compute **precision** for n-grams:
 - Unigrams -> BLEU-1
 - Bigrams -> BLEU-2
 - Trigrams -> BLEU-3
 - 4-grams -> BLEU-4

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Precision and Recall

$$\text{Precision} = \frac{\text{Relevant and retrieved instances}}{\text{All retrieved instances}} \leftarrow \text{Predicted by a model}$$

$$\text{Recall} = \frac{\text{Relevant and retrieved instances}}{\text{All relevant instances}} \leftarrow \text{Ground-truths}$$

Relevant and retrieved instances: **Intersection** between predictions and ground-truths

Calculation of BLEU Score (Example)

Assume we now translate from Chinese to English.

Calculate BLEU-1 score

Chinese: 我想要讀那本書

Reference1: I want to read the book.

Reference2: I want to read that book.

Model output: the the the the the the.

Calculation of BLEU Score (Example)

Assume we now translate from Chinese to English.

Calculate BLEU-1 score

Chinese: 我想要讀那本書

Precision: $\frac{6}{6}$

Reference1: I want to read the book.

Reference2: I want to read that book.

100%! Can this be true?

Model output: the the the the the the.

Calculation of BLEU Score (Example)

Assume we now translate from Chinese to English.

Calculate BLEU-1 score

Chinese: 我想要讀那本書

~~Precision: $\frac{6}{6}$~~

Reference1: I want to read the book.

Reference2: I want to read that book.

Modified Precision: $\frac{1}{6}$

Model output: the the the the the the the.

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Why should we use modified precision?

- The output sequences can be total mistakes.
 - E.g., the the the the the the
- Original precision is in favor of **longer** output sequences.
- Therefore, we should use modified precision to prevent bad evaluations.

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Calculation of BLEU Score (Example)

Calculate BLEU-2 score

Reference1: The dog is on the bed.

Reference2: There is a dog on the bed.

Model output: The dog the dog on the bed.

← More than one references can be provided for machine translation!

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation." Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Calculation of BLEU Score (Example)

Calculate BLEU-2 score

		Count	
Reference1: The dog is on the bed.	the dog	2	(duplicated)
Reference2: There is a dog on the bed.	dog the	1	
Model output: <u>The dog</u> the dog <u>on the</u> bed.	dog on	1	
<u>1</u> <u>2</u> <u>3</u> <u>4</u> <u>5</u> <u>6</u>	on the	1	
	the bed	1	

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Calculation of BLEU Score (Example)

Calculate BLEU-2 score

Reference1: The dog is on the bed.

Reference2: There is a dog on the bed.

Model output: The dog the dog on the bed.

	Count	Clips to the reference ↓ Count _{clip}
the dog	2	1
dog the	1	
dog on	1	
on the	1	
the bed	1	

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Calculation of BLEU Score (Example)

Calculate BLEU-2 score

	Count	Count _{clip}
Reference1: The dog is on the bed.	the dog 2	1
Reference2: There is a dog on the bed.	dog the 1	0
Model output: The <u>dog the</u> dog on the bed.	dog on 1	
	on the 1	
	the bed 1	

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Calculation of BLEU Score (Example)

Calculate BLEU-2 score

	Count	Count _{clip}
Reference1: The dog is on the bed.	the dog 2	1
Reference2: There is a <u>dog on</u> the bed.	dog the 1	0
Model output: The dog the <u>dog on</u> the bed.	dog on 1	1
	on the 1	
	the bed 1	

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Calculation of BLEU Score (Example)

Calculate BLEU-2 score

	Count	Count _{clip}
Reference1: The dog is <u>on the</u> bed.	the dog 2	1
Reference2: There is a dog <u>on the</u> bed.	dog the 1	0
Model output: The dog the dog <u>on the</u> bed.	dog on 1	1
Count only one time even mapped to both references.	on the 1	1
	the bed 1	

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Calculation of BLEU Score (Example)

Calculate BLEU-2 score

	Count	Count _{clip}
Reference1: The dog is on <u>the bed</u> .	the dog 2	1
Reference2: There is a dog on <u>the bed</u> .	dog the 1	0
Model output: The dog the dog on <u>the bed</u> .	dog on 1	1
	on the 1	1
	the bed 1	1

Count **only one time** even mapped to both references.

Modified Precision: $\frac{4}{6}$

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Formula of BLEU Score (1)

Summation for unigram, bigram, tri-gram, and 4-gram

$$p_n = \frac{\sum_{C \in \{Candidates\}} \sum_{n\text{-gram} \in C} Count_{clip}(n\text{-gram})}{\sum_{C' \in \{Candidates\}} \sum_{n\text{-gram}' \in C'} Count(n\text{-gram}')}$$

Summation for all candidates (model outputs)
of each translation

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Formula of BLEU Score (2)

Summation for unigram, bigram, tri-gram, and 4-gram

$$p_n = \frac{\sum_{C \in \{Candidates\}} \sum_{n\text{-gram} \in C} Count_{clip}(n\text{-gram})}{\sum_{C' \in \{Candidates\}} \sum_{n\text{-gram}' \in C'} Count(n\text{-gram}')}$$

Summation for all candidates (model outputs)
of each translation

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

What we've learned BLEU so far

- The BLEU score is calculated from the summation of 1-gram to 4-gram.
 - You can also measure n-gram individually.
- We use modified precision to prevent bad evaluations.
- What will happen if a model tends to generate really short sentences?



More penalty for calculating BLEU score!

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

Brevity Penalty (BP)

- BP is used to penalize **short** candidates.

c : The length of a candidate sequence
 r : The length of a reference sequence that is closest to c (shorter one)

$$\text{BP} = \begin{cases} 1 & \text{if } c > r \\ e^{(1-r/c)} & \text{if } c \leq r \end{cases}$$

Then,

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{n=1}^N w_n \log p_n \right)$$

$N=4$ to include 1-gram to 4-gram

Weight for each n -gram (was set 1/4 in the original paper)

Papineni, Kishore, et al. "Bleu: a method for automatic evaluation of machine translation."
Proceedings of the 40th annual meeting of the Association for Computational Linguistics. 2002.

ROUGE Score

- ROUGE (Recall-Oriented Understudy for Gisting Evaluation)
- Mainly for text summarization
- Metric Input: **Summary** (prediction), **Reference** (gold summary)
- Common metrics: ROUGE-1, ROUGE-2, ROUGE-L
 - L: Longest common subsequence
- **Please note that current papers calculate ROUGE-F as default!!!**
 - In other words, ROUGE-1**F**, ROUGE-2**F**, ROUGE-L**F**

Lin, Chin-Yew. "Rouge: A package for automatic evaluation of summaries." Text summarization branches out. 2004.

ROUGE-1 Example

predictions = ["The", "cat", "sat", "on", "the", "mat"]

references = ["A", "cat", "was", "sitting", "on", "the", "mat"]

ROUGE-1 **recall** = Number of matching unigrams / Number of unigrams in the reference = $4/7$

ROUGE-1 **precision** = Number of matching unigrams / Number of unigrams in the machine-generated summary = $4/6$

ROUGE-1 **F1-score** = Harmonic mean of the precision and the recall = $2 * 4/7 * 4/6 / (4/7 + 4/6)$

ROUGE-2 Example

predictions = ["The cat", "cat sat", "sat on", "on the", "the mat"]

references = ["A cat", "cat was", "was sitting", "sitting on", "on the", "the mat"]

ROUGE-2 **recall** = Number of matching bigrams / Number of bigrams in the reference = 2/6

ROUGE-2 **precision** = Number of matching bigrams / Number of bigrams in the machine-generated summary = 2/5

ROUGE-2 **F1-score** = Harmonic mean of the precision and the recall = $2 * 2/6 * 2/5 / (2/6 + 2/5)$

ROUGE-L Example

predictions = ["The", "cat", "sat", "on", "the", "mat"]

references = ["A", "cat", "was", "sitting", "on", "the", "mat"]

The longest common subsequence is ["cat", "on", "the", "mat"]

ROUGE-L **recall** = Number of matching unigrams / Number of unigrams in the reference = $4/7$

ROUGE-L **precision** = Number of matching unigrams / Number of unigrams in the machine-generated summary = $4/6$

ROUGE-L **F1-score** = Harmonic mean of the precision and the recall = $2 * 4/7 * 4/6 / (4/7 + 4/6)$

ROUGE-L Example

The order should be kept for the LCS problem

predictions = ["The", "cat", "sat", "on", "the", "mat"]

references = ["on", "the", "mat", "sitting", "a", "cat"]

The longest common subsequence is ["on", "the", "mat"]

ROUGE-L **recall** = Number of matching unigrams / Number of unigrams in the reference = 3/6

ROUGE-L **precision** = Number of matching unigrams / Number of unigrams in the machine-generated summary = 3/6

ROUGE-L **F1-score** = Harmonic mean of the precision and the recall = $2 * 0.5 * 0.5 / (0.5 + 0.5)$

Why do we need BLEU and ROUGE?

- BLEU is mainly designed for machine translation.
 - For example, the **Brevity Penalty**.
- ROUGE measures the overlapping between predicted and gold summaries.
- Can we just use one of them?
 - Conventionally, no.
 - Different tasks are evaluated with different metrics.

Comparison for Human and Automatic Evaluations (e.g., BLEU and ROUGE)

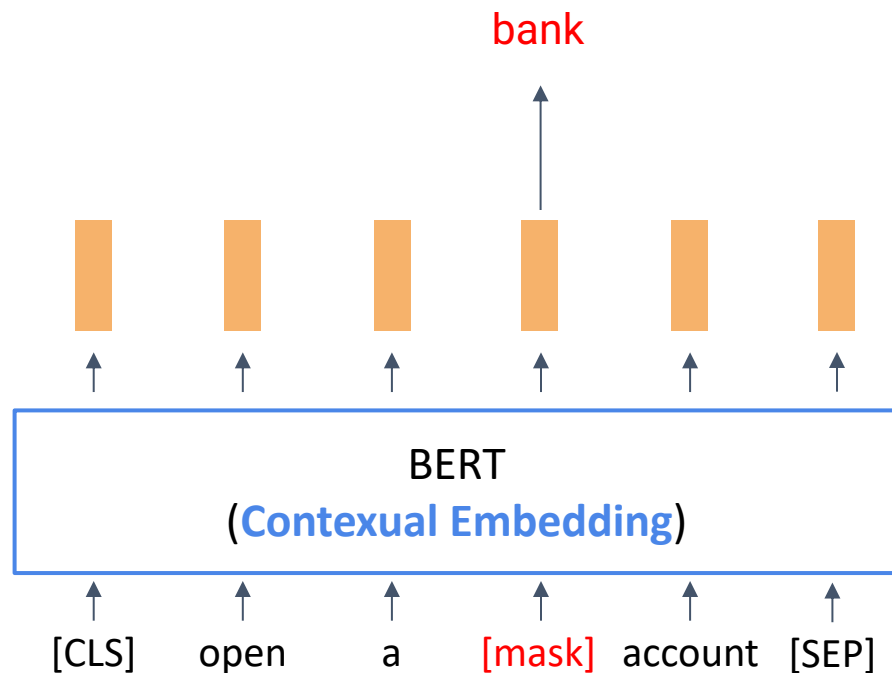
- **Human evaluations**
 - Pros: More accurate for subjectivity, flexibility for any desired comparison
 - Cons: Less objective, time-consuming, expensive
- Automatic evaluations
 - Pros: Objective enough to serve as common evaluation metrics, fast
 - Cons: Cannot meet language diversity
 - Take machine translation for instance, there are always other valid ways to translate the source sentence.

Issue of BLEU and ROUGE


- Cons: Cannot meet language diversity
 - This mainly comes from the way for measuring **overlapping** rates.
- **Question:** **Can we create an automatic metric to fix the issue?**
- Next, we are going to introduce two **learned** automatic evaluation metrics
 - **BERTScore** (ICLR 2020)
 - **BLEURT** (ACL 2020)

(Recap) BERT: Bidirectional Encoder Representations from Transformers

Masked Language Modelling




Next Sentence Prediction



Devlin, Jacob, et al. "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding." NAACL 2019.

BERTScore – Overview



Zhang, Tianyi, et al. "BERTScore: Evaluating Text Generation with BERT." International Conference on Learning Representations. 2020.

BERTScore – Steps

Step 0: Prepare Reference x , Candidate \hat{x} , and a pre-trained BERT model

Step 1: Infer x and \hat{x} with BERT respectively, get a sequence of output vectors

$\langle \mathbf{x}_1, \dots, \mathbf{x}_k \rangle$ for x and a sequence of output vectors $\langle \hat{\mathbf{x}}_1, \dots, \hat{\mathbf{x}}_k \rangle$ for \hat{x}

Zhang, Tianyi, et al. "BERTScore: Evaluating Text Generation with BERT." International Conference on Learning Representations. 2020.

BERTScore – Steps

Zhang, Tianyi, et al. "BERTScore: Evaluating Text Generation with BERT." International Conference on Learning Representations. 2020.

Step 0: Prepare Reference x , Candidate \hat{x} , and a pre-trained BERT model

Step 1: Infer x and \hat{x} with BERT respectively, get a sequence of output vectors

$\langle x_1, \dots, x_k \rangle$ for x and a sequence of output vectors $\langle \hat{x}_1, \dots, \hat{x}_k \rangle$ for \hat{x}

Step 2: Measure pairwise cosine similarity

Recall

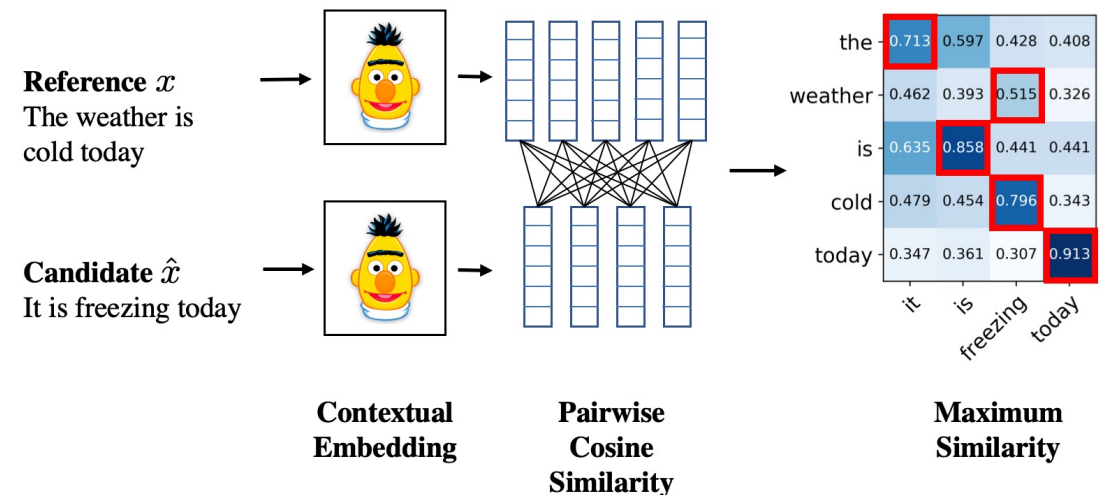
$$R_{\text{BERT}} = \frac{1}{|x|} \sum_{x_i \in x} \max_{\hat{x}_j \in \hat{x}} x_i^\top \hat{x}_j$$

Based on reference

Precision

$$P_{\text{BERT}} = \frac{1}{|\hat{x}|} \sum_{\hat{x}_j \in \hat{x}} \max_{x_i \in x} x_i^\top \hat{x}_j$$

Based on candidate



BERTScore – Importance Weighting

Given M reference sentences $\{x^{(i)}\}_{i=1}^M$, the idf (inverse document frequency) score of a word-piece token w is:

$$\text{Idf}(w) = -\log \frac{1}{M} \sum_{i=1}^M \mathbb{I}[w \in x^{(i)}]$$

$$R_{\text{BERT}} = \frac{\sum_{x_i \in x} \text{idf}(x_i) \max_{\hat{x}_j \in \hat{x}} \mathbf{x}_i^\top \hat{\mathbf{x}}_j}{\sum_{x_i \in x} \text{idf}(x_i)}$$

the	0.713	0.597	0.428	0.408	1.27
weather	0.462	0.393	0.515	0.326	7.94
is	0.635	0.858	0.441	0.441	1.82
cold	0.479	0.454	0.796	0.343	7.90
today	0.347	0.361	0.307	0.913	8.88
	it	is	freezing	today	idf weights

**Maximum
Similarity**

Importance Weighting

$$\rightarrow R_{\text{BERT}} = \frac{(0.713 \times 1.27) + (0.515 \times 7.94) + \dots}{1.27 + 7.94 + 1.82 + 7.90 + 8.88} = 0.753$$

Zhang, Tianyi, et al. "BERTScore: Evaluating Text Generation with BERT." International Conference on Learning Representations. 2020.

Summary of BERTScore

- BERTScore leverages the contextual representation abilities of BERT to measure the semantic similarities between a reference and a candidate.
- In the paper, BERTScore correlates better with human judgments and provides stronger model selection performance than existing metrics.
- However, BERTScore does not involve training process.

Can we train BERT for a better evaluation metric?

Zhang, Tianyi, et al. "BERTScore: Evaluating Text Generation with BERT." International Conference on Learning Representations. 2020.

BLEURT – Quick Introduction

- BLEURT: Learning Robust Metrics for Text Generation, published by Google
- BLEURT **trains** BERT for a more robust evaluation metric.
 - Mainly for **machine translation**.
 - Also get hints from the name “BLEURT”
 - Trained checkpoint can be obtained. We don't need to perform training.


Sellam, Thibault, Dipanjan Das, and Ankur Parikh. "BLEURT: Learning Robust Metrics for Text Generation."
Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics. 2020.

BLEURT – Motivations

- **Learned** metrics can be tuned to **measure task-specific** properties, such as fluency, faithfulness, grammar, or style.
- NLG systems tend to get better over time, and therefore a model trained on ratings data from 2015 may fail to distinguish top performing systems in 2019, especially for newer research tasks.


Training on Human Ratings

Traditional




Learning

Training



Inference



BLEURT – Steps

Step 0: Reference-candidate pairs (z, \tilde{z}) and the pre-trained BERT model

Step 1: Data augmentation for (z, \tilde{z}) to perform pre-training

Data augmentation strategies

- Random masking
- Back-translation
- Dropping words randomly

Total 6.5 million variants of (z, \tilde{z}) were created.

Sellam, Thibault, Dipanjan Das, and Ankur Parikh. "BLEURT: Learning Robust Metrics for Text Generation."
Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics. 2020.

Random masking

Two kinds of masking strategies were adopted:

Token masking

I love traveling to Vancouver for attending a conference.



I love traveling to Vancouver for **[MASK]** a conference.

Span masking

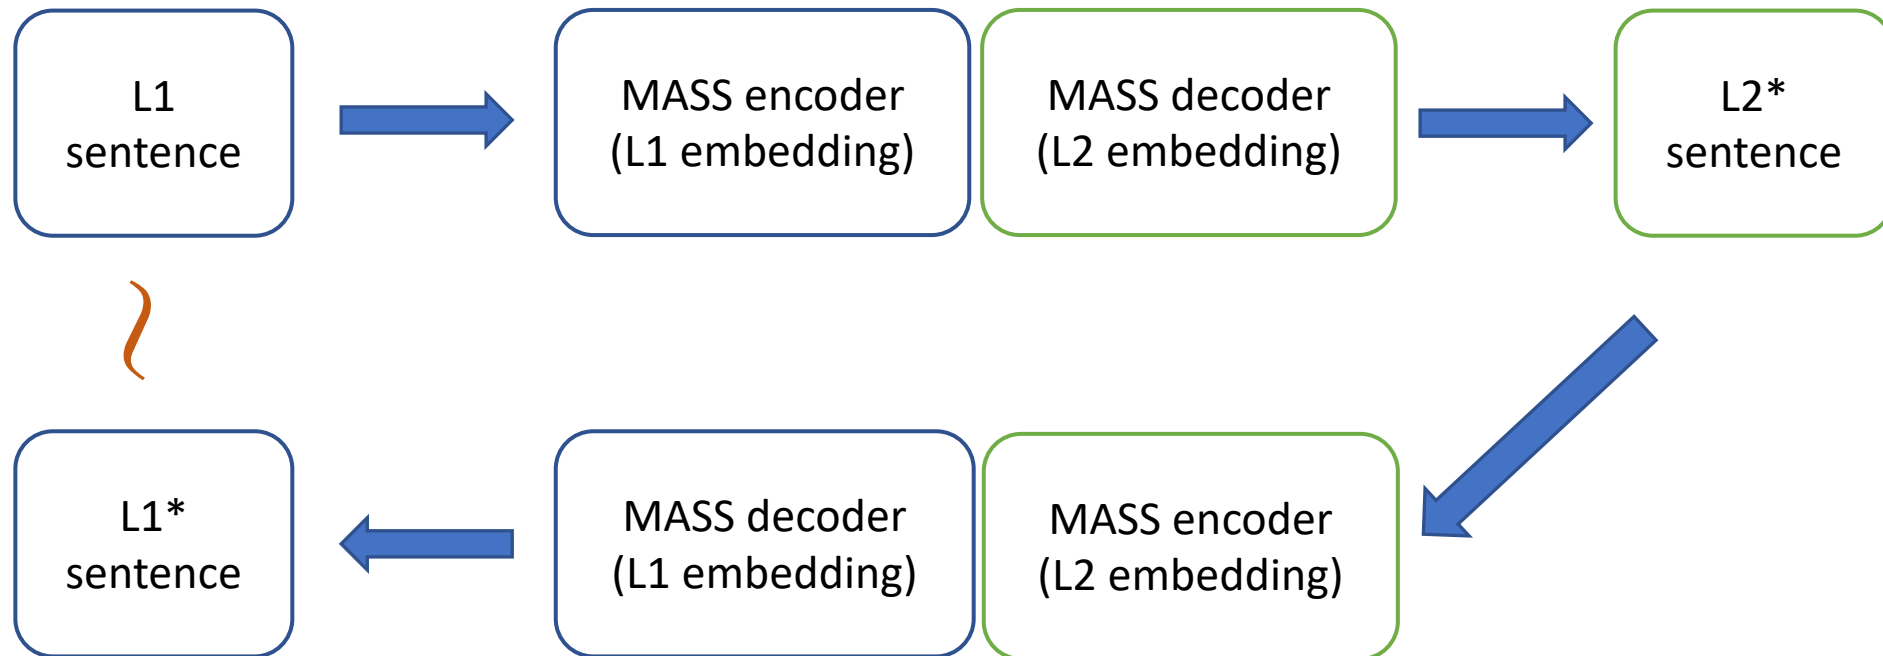
I love traveling to Vancouver for attending a conference.



I love traveling to Vancouver for **[MASK] [MASK] [MASK]**.

Backtranslation

- L1: English; L2: French or German



Dropping words randomly

- The authors found it useful in their experiments to randomly drop words to create other examples.

I love traveling to Vancouver for attending a conference.



I love to Vancouver for attending a conference.

Sellam, Thibault, Dipanjan Das, and Ankur Parikh. "BLEURT: Learning Robust Metrics for Text Generation." Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics. 2020.

BLEURT – Step 3

Step 3: Pre-training each sentence pair (z, \tilde{z}) with the following tasks.

Note that this is not conventional BERT pre-training! It is **multi-task pre-training**!

Task Type	Pre-training Signals	Loss Type
BLEU	τ_{BLEU}	Regression
ROUGE	$\tau_{\text{ROUGE}} = (\tau_{\text{ROUGE-P}}, \tau_{\text{ROUGE-R}}, \tau_{\text{ROUGE-F}})$	Regression
BERTscore	$\tau_{\text{BERTscore}} = (\tau_{\text{BERTscore-P}}, \tau_{\text{BERTscore-R}}, \tau_{\text{BERTscore-F}})$	Regression
Backtrans. likelihood	$\tau_{\text{en-fr}, z \tilde{z}}, \tau_{\text{en-fr}, \tilde{z} z}, \tau_{\text{en-de}, z \tilde{z}}, \tau_{\text{en-de}, \tilde{z} z}$	Regression
Entailment	$\tau_{\text{entail}} = (\tau_{\text{Entail}}, \tau_{\text{Contradict}}, \tau_{\text{Neutral}})$	Multiclass
Backtrans. flag	$\tau_{\text{backtran_flag}}$	Multiclass

- Ground-truth values can be computed for each (z, \tilde{z}) pair!
- Losses for the six tasks were sum up during pre-training.
- Regression: mean squared error
- Multiclass: Cross-entropy

Sellam, Thibault, Dipanjan Das, and Ankur Parikh. "BLEURT: Learning Robust Metrics for Text Generation." Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics. 2020.

Task 4: Backtranslation Likelihood

- Existing translation models (trained) are needed.
 - Transformers (Vaswani et al., 2017): EN-FR
 - Transformers (Vaswani et al., 2017): DE-EN
- Equations use EN-FR for an example

$$z_{\text{fr}}^* = \arg \max P_{\text{en} \rightarrow \text{fr}}(z_{\text{fr}} | z)$$

Best translated French sentence (details absent in the paper)

$$P(\tilde{z} | z) \approx P_{\text{fr} \rightarrow \text{en}}(\tilde{z} | z_{\text{fr}}^*)$$

Backtranslation Likelihood

$P(x_t | x_1, \dots, x_{t-1}, z)$

Task 5 and Task 6

- **Textual Entailment**
 - We report the probability of three labels: **Entail**, **Contradict**, and **Neutral**, using BERT fine-tuned on the MNLI dataset.
- **Backtranslation flag**
 - A **Boolean** that indicates whether the perturbation was generated with backtranslation or with mask-filling.

BLEURT – Final Step

Step 4: Fine-tune the model (trained from Step 3) on the **<Reference, Candidate ,Rating> data** using the regression loss

The **<Reference, Candidate ,Rating> data** include

- WMT (machine translation task)
- WebNLG (for general text generation)

Summary of BLEURT

- This approach uses (continual) pre-training and fine-tuning to create a learned evaluation metric for machine translation and general NLG.
- According to the paper, BLEURT is better aligned to human ratings than BERTScore.
- BLEURT should work for text summarization, but the authors did not test it.

Sellam, Thibault, Dipanjan Das, and Ankur Parikh. "BLEURT: Learning Robust Metrics for Text Generation." Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics. 2020.

Thank you!

Instructor: 林英嘉

 yjlin@cgu.edu.tw