

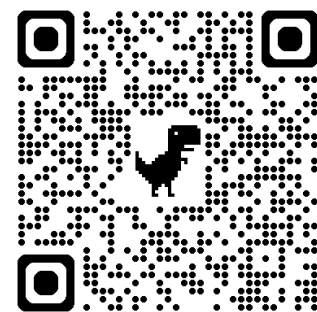


自然語言處理與應用

Natural Language Processing and Applications

如何訓練大模型？輕量化微調方法
Parameter-efficient Fine-tuning (PEFT)

Instructor: 林英嘉 (Ying-Jia Lin)
2025/05/05



[Course GitHub](#)



[Slido # NLP_0505](#)

The Revolution of ChatGPT

ChatGPT came out in November, 2022.

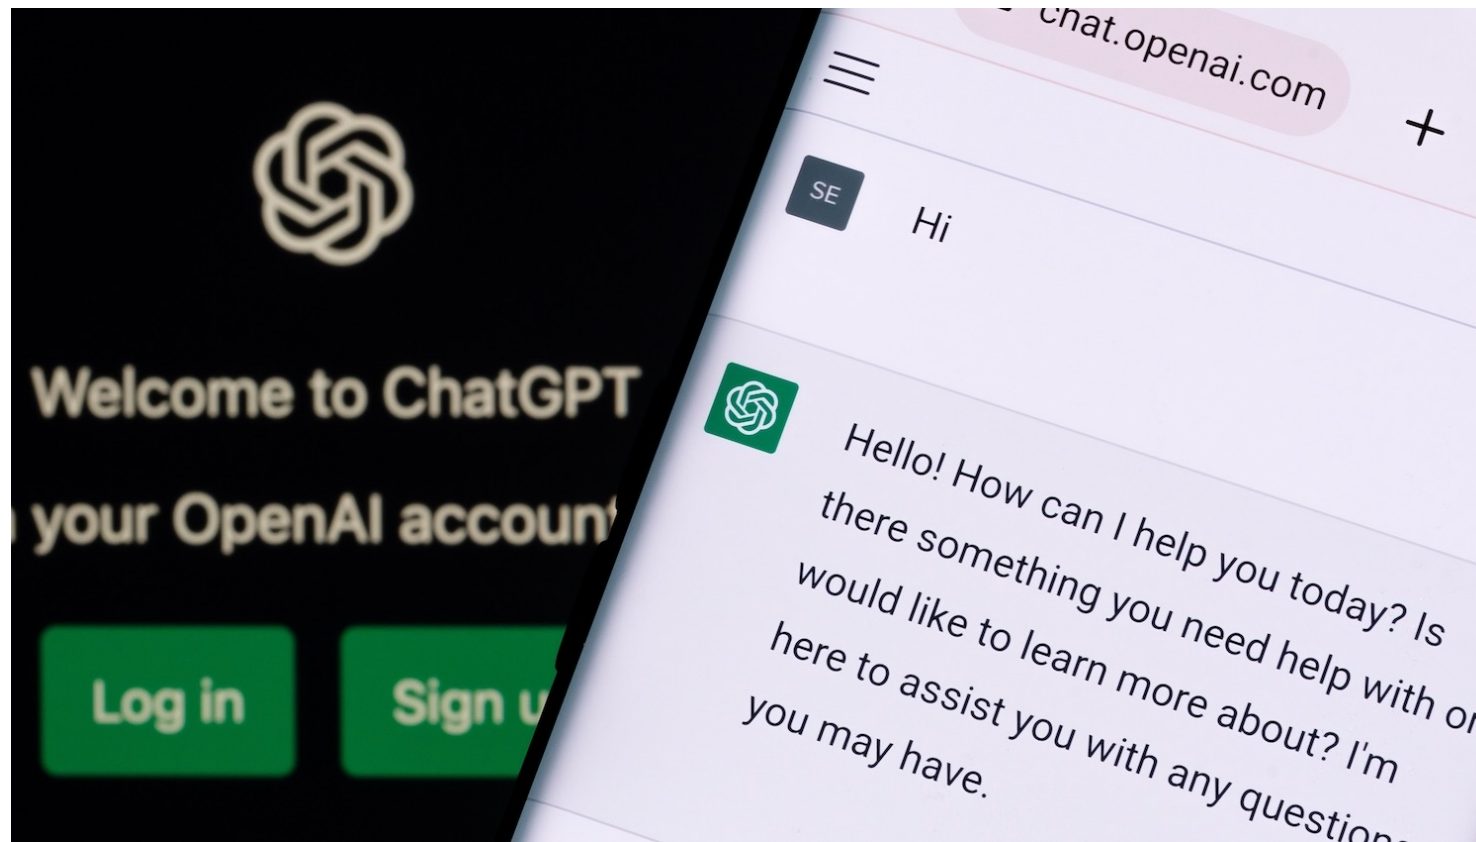



Figure source: <https://poole.ncsu.edu/thought-leadership/article/lets-chat-about-chatgpt/>
<https://openai.com/index/chatgpt/>

讓大型語言模型適用於你的任務？



翻譯、聊天、寫故事 ...



醫學資料、少數語言、網路上查不到的知識




解決方案：Fine-tuning 語言模型

Figure source:
https://gameofthrones.fandom.com/wiki/High_Valyrian?file=Nekeesse_Valyrio.jpg

Full Fine-tuning (全微調) LLM 的困境

LARGE LANGUAGE MODEL HIGHLIGHTS (FEB/2024)



Sizes linear to scale. Selected highlights only. All models are available. All models are Chinchilla-aligned (20:1 tokens:parameters) <https://lilearchitect.ai/chinchilla/> All 200+ models: <https://lilearchitect.ai/models-table/> Alan D. Thompson, 2023-2024.

Source: Inside language models (from GPT-4 to PaLM) – Dr Alan D. Thompson – Life Architect

GPU Memory Estimated (Model weights)

以 Llama 2-7B 為例，且儲存參數的精度為 float 32 (FP32)：

在 float 32 (FP32) 的情況下，1 個參數需要 4 bytes 儲存

-> 7B = 70億個參數 = $7,000,000,000 * 4 \text{ bytes}$ = (扣掉3個逗號從B換成GB) **28GB**

Introduction to Quantization

(這頁只是示意圖，數值不精準)

$$\begin{bmatrix} -0.4 & 1.3 & 3.73 \\ -4.7 & -3.2 & -6.4 \\ 8.5 & 14.3 & 13.5 \end{bmatrix} \xrightarrow[\text{FP32} \rightarrow \text{int8}]{\text{降低儲存精度}} \begin{bmatrix} 0 & 1 & 4 \\ -5 & -3 & -6 \\ 9 & 14 & 14 \end{bmatrix}$$

36 bytes **8 bytes**

32-bit floating point (FP32): 1個
值需要**4個bytes**才能儲存

8-bit Integer (int8): 1個值需要
1個bytes才能儲存

誤差：

$$\begin{bmatrix} 0.4 & -0.3 & 0.27 \\ -0.3 & 0.2 & 0.4 \\ 0.5 & -0.3 & 0.5 \end{bmatrix}$$

GPU Memory Estimated (Model weights)

以 Llama 2-7B 為例，且儲存參數的精度為 float 32 (FP32)：

在 float 32 (FP32) 的情況下，1 個參數需要 4 bytes 儲存

-> 7B = 70億個參數 = $7,000,000,000 * 4 \text{ bytes}$ = (扣掉3個逗號從B換成GB) **28GB**

模型	參數量	Memory (FP32)	Memory (FP16)
DeepSeek v3	685B	2740 GB	1370 GB
Llama 4 Scout	109B	436 GB	218 GB
GPT-2 XL	1.5B	6 GB	3 GB


GPU Memory Estimated (Full Fine-tuning)

Llama 2-7B		
16-bit float, max_length (seq) = 4096, hidden_size = 4096, batch_size (bs) = 1		
	算法	Memory
CUDA	-	~1 GB
Model weights	$\text{size(float)} * N_{\text{parameter}}$	13.03 GB
Gradients	$\text{size(float32)} * N_{\text{trainable}}$	26.06 GB
Hidden states	$\sim \text{size(float)} * \text{seq} * \text{hidden_size} * L$	1.07 GB
Optimizer states (Adam)	$2 * \text{size(float)} * N_{\text{trainable}}$	26.06 GB


L : Number of layers in model (eq. 32 layers)

Estimate: 67.22 GB
*NVIDIA 5090: 32GB

有沒有可能只微調 LLM 一部分的參數呢？



Full Fine-tuning
Update **all model parameters**



Parameter-efficient Fine-tuning (PEFT)
Update a **small subset** of model parameters

GPU Memory Estimated (PEFT)

- 假設我們只更新 2% 的參數

Llama 2-7B 16-bit float, max_length (seq) = 4096, hidden_size = 4096, batch_size (bs) = 1		
	算法	Memory
CUDA	-	~1 GB
Model weights	$\text{size(float)} * N_{\text{parameter}}$	13.03 GB
Gradients	$\text{size(float32)} * N_{\text{trainable}}$	$13.03 * 0.02 = 0.2606 \text{ GB}$
Hidden states	$\sim \text{size(float)} * \text{seq} * \text{hidden_size} * L$	1.07 GB
Optimizer states (Adam)	$2 * \text{size(float)} * N_{\text{trainable}}$	0.5212 GB

L : Number of layers in model (eq. 32 layers)

Estimate: **15.88 GB**

*NVIDIA 4060 Ti: 16 GB

PEFT Outline

Adapters

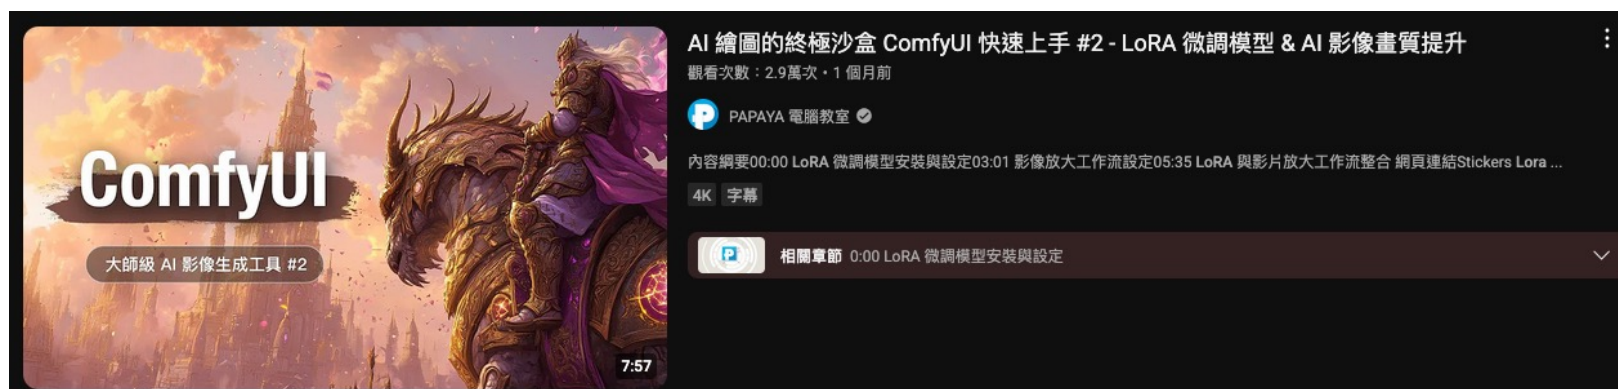
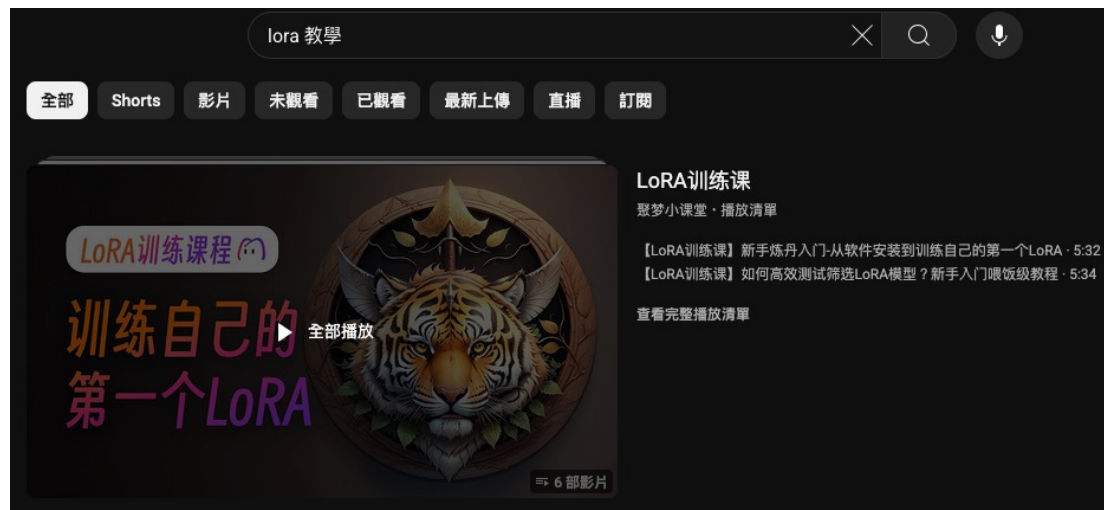
LoRA

Prefix-
Tuning


BitFit

LoRA: **Lo**w-**R**ank **A**daptation

LoRA is common ...




[Prerequisite] Weights in Transformer layers




$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

[Recap] MLP is composed of weight matrices



$$\begin{matrix} [x_1 & x_2 & x_3] \\ \mathbb{R}^{1 \times 3} \end{matrix} \times \begin{matrix} \begin{bmatrix} w_{1,1} & w_{1,2} \\ w_{2,1} & w_{2,2} \\ w_{3,1} & w_{3,2} \end{bmatrix} \\ \mathbb{R}^{3 \times 2} \end{matrix} + \text{bias} \quad \mathbb{R}^2$$

[Prerequisite] 矩陣分解




參數量 $1000 \times 1000 = 1\text{百萬}$

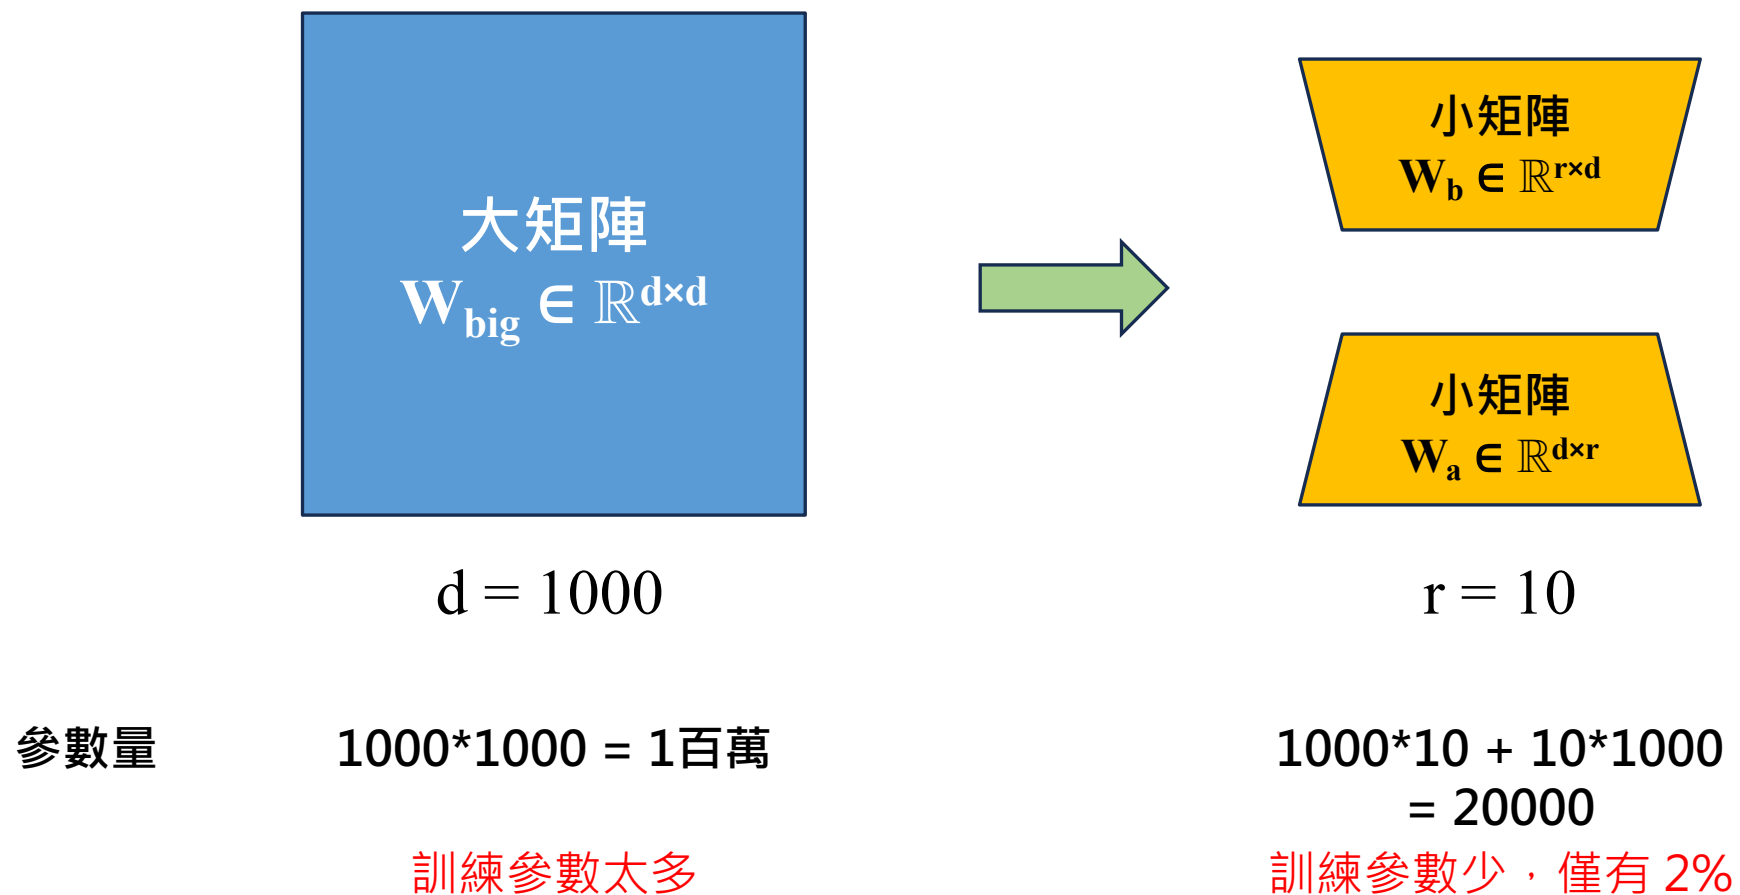
$1000 \times 10 + 10 \times 1000$
 $= 20000$

LoRA: Low-Rank Adaptation

只有黃色部分會被訓練
 α 是超參數



為什麼 LoRA 要這樣做？



Low-Rank 的部分在哪？(1)


$$\begin{bmatrix} \checkmark & \checkmark & & \checkmark \\ 1 & 0 & -1 & 0 \\ 0 & 1 & 2 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Reduced row-
echelon form

\checkmark : Linearly-independent vectors
rank = 3

rank 大小最多等於 column vectors
(或 row vectors) 的數量

Low-Rank 的部分在哪？(2)



$d > r$ 的情況下 (通常 $d \gg r$)，**A 或 B** 的 rank 一定遠小於 **W** 的 rank
故為 low-rank 的由來

LoRA: Low-Rank Adaptation (w / pseudo code)

Pseudocode:

```
input_dim = 768 # the hidden size of the pre-trained model
output_dim = 768 # the output size of the layer
rank = 8 # The rank 'r' for the low-rank adaptation

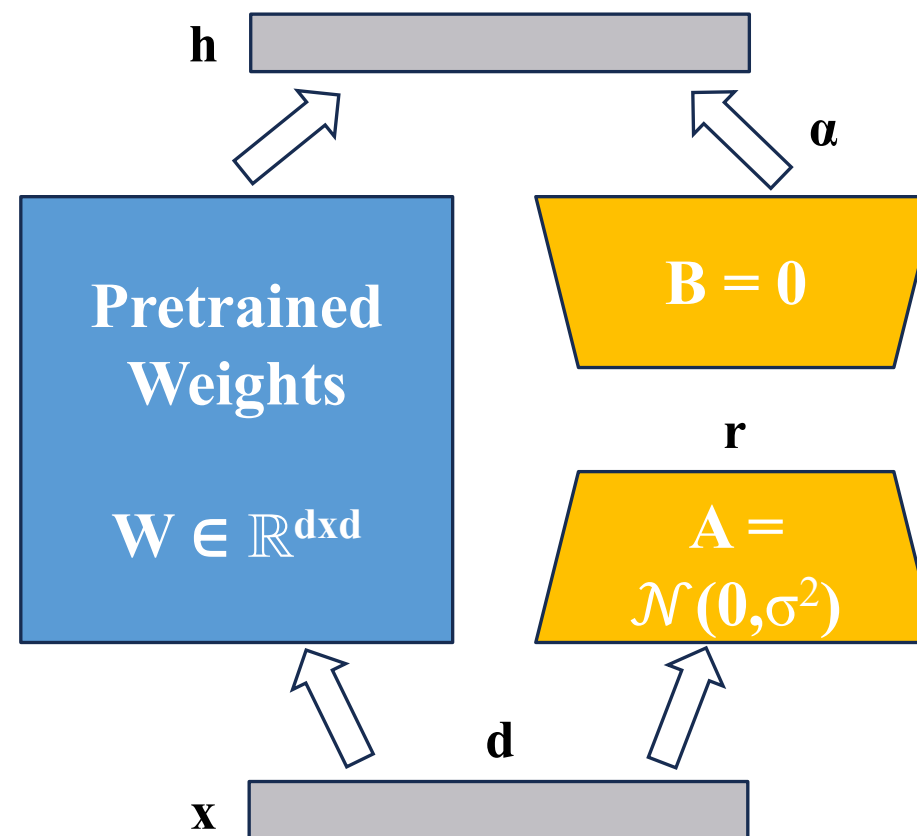
W = ... # from pretrained network with shape input_dim x
output_dim

W_A = nn.Parameter(torch.empty(input_dim, rank)) # LoRA weight A
W_B = nn.Parameter(torch.empty(rank, output_dim)) # LoRA weight B

# Initialization of LoRA weights
nn.init.kaiming_uniform_(W_A, a=math.sqrt(5))
nn.init.zeros_(W_B)


def regular_forward_matmul(x, W):
    h = x @ W
    return h

def lora_forward_matmul(x, W, W_A, W_B):
    h = x @ W # regular matrix multiplication
    h += x @ (W_A @ W_B) * alpha # use scaled LoRA weights
    return h
```



[注意事項] LoRA: Low-Rank Adaptation

- r 的數值大小需要手動調整
 - r 越小，訓練參數量越少
- 相較於原本沒有 LoRA 的模型，
LoRA 其實會讓 inference 速度變慢



How about model compression?



PEFT vs. Model compression

	PEFT	Model Compression
目標	讓模型適應新的任務， 但模型大小不變	加速模型運算或模型儲存空間
相較於原始模型的 改變內容	插入少量可訓練參數	減少整體模型結構或權重
參數更新	只更新少量新參數	先減少整體模型結構後針對新的模 型進行訓練
使用情境	需要模型學會新的任務時	手機、邊緣裝置

Thank you!

Instructor: 林英嘉

 yjlin@cgu.edu.tw

TA: 吳宣毅

 m1161007@cgu.edu.tw